



Haswell-EP Server Reference Code User's Guide

March 28th, 2014
Revision 0.96

INTEL CONFIDENTIAL - NDA

Copyright 2006 – 2014 Intel Corporation All Rights Reserved. This document contains trade secrets and proprietary information of Intel Corporation or its suppliers and licensors, and is protected by worldwide copyright and trade secret laws and treaty provisions. No part of the Material may be used, copied, reproduced, modified, published, uploaded, posted, transmitted, distributed, or disclosed in any way without Intel's prior express written permission.



INTEL CONFIDENTIAL - NDA

Copyright 2006 – 2014 Intel Corporation All Rights Reserved.

The source code contained or described herein and all documents related to the source code ("Material") are owned by Intel Corporation or its suppliers. Title to the Material remains with Intel Corporation, its suppliers, or licensors. The Material contains trade secrets and proprietary and confidential information of Intel Corporation, its suppliers, and licensors, and is protected by worldwide copyright and trade secret laws and treaty provisions. No part of the Material may be used, copied, reproduced, modified, published, uploaded, posted, transmitted, distributed, or disclosed in any way without Intel's prior express written permission.

No license under any patent, copyright, trade secret or other intellectual property right is granted to or conferred upon you by disclosure or delivery of the Materials, either expressly, by implication, inducement, estoppel or otherwise. Any license under such intellectual property rights must be express and approved by Intel in writing.

Unless otherwise agreed by Intel in writing, you may not remove or alter this notice or any other notice embedded in Materials by Intel or Intel's suppliers or licensors in any way.

Contents

1	INTRODUCTION.....	7
1.1	REFERENCE CODE PACKAGE CONTENTS.....	7
2	MINIBIOS OVERVIEW	8
2.1	MINIBIOS	8
2.2	BUILD TOOLS.....	8
2.3	BUILD OPTIONS	8
2.4	DIRECTORY STRUCTURE.....	8
2.5	BUILD ENVIRONMENT.....	10
2.6	BUILDING THE MINIBIOS	10
2.7	DEBUGGING THE MINIBIOS USING INTEL SOFTSIMULATION	11
2.7.1	Configuring SoftSimulation to Load the BIOS image file	11
2.7.2	Symbolic Debugging of MiniBIOS	11
2.7.3	Using the DUMPBIN utility	11
3	DETAILS OF MINIBIOS.....	14
3.1	MINIBIOS LAYOUT	14
	FIGURE 1: MINIBIOS LAYOUT	14
3.2	BUILDING THE MINIBIOS	14
3.3	EXECUTION FLOW	15
3.4	MINIBIOS TOOLS	15
3.4.1	FindEntryPoint.....	15
3.4.2	PlacePE	15
4	COMMON REFERENCE CODE COMPONENTS	17
4.1	OVERVIEW.....	17
4.2	BIOS REFERENCE CODE HIGH LEVEL EXECUTION FLOW	17
4.3	ASSUMPTIONS	17
4.4	PLATFORM PORTING AND REFERENCE CODE HOOKS	17
4.4.1	Options.....	17
4.4.2	lowGap.....	18
4.4.3	highGap.....	18
4.4.4	serialDebugMsgLvl.....	18
4.4.5	bsdBreakpoint	18
4.4.6	maxAddrMem.....	18
4.4.7	SocketType	18
4.4.8	debugPort.....	18
4.4.9	nvrampPtr.....	18
4.4.10	sysHostBufferPtr	18
4.4.11	mmCfgBase	18
4.4.12	mmCfgSize.....	18
4.4.13	mmiolBase	19
4.4.14	mmiolSize.....	19
4.4.15	mmiohBase	19
4.4.16	mmiohSize	19
4.4.17	pchUmaEn	19
4.4.18	numaEn.....	19
4.4.19	isocEn	19
4.4.20	DciVcm.....	19
4.4.21	MeSegEn	19
4.4.22	dcaEn.....	19
4.4.23	PSMIEncabledSupport.....	19
4.4.24	BdatEn	19
4.4.25	consoleComPort.....	19
4.4.26	debugComPort.....	19
4.4.27	usbDebugPortSetup.....	19
4.5	SCRATCH PAD REGISTER USAGE	20
4.6	HIGH AND LOW MMIO BAR ALLOCATION	20
5	CPU INITIALIZATION REFERENCE CODE	20



5.1	OVERVIEW.....	20
5.2	CPU REFERENCE CODE EXECUTION FLOW	20
5.3	ASSUMPTIONS	21
5.4	PLATFORM BIOS INTEGRATION GUIDE	21
5.4.1	Build-Time Equates and Defines	21
5.4.2	Input Parameters.....	22
5.5	ROUTINES FOR PLATFORM CUSTOMIZATION	22
5.6	POST CODES AND SERIAL PORT DEBUG MESSAGES	23
5.6.1	POST Code at IO Port 0x80.....	23
5.6.2	Serial Port Debug Messages.....	23
6	QPI REFERENCE CODE.....	25
6.1	OVERVIEW.....	25
6.2	INPUT PARAMETERS	25
6.2.1	PPINrOptIn.....	25
6.2.2	BusRatio/IoRatio/MmioIRatio	25
6.2.3	LegacyVgaSoc	25
6.2.4	MmioP2pDis.....	25
6.2.5	IsocAzaliaVc1En	26
6.2.6	DebugPrintLevel.....	26
6.2.7	AltRtid2S	26
6.2.8	ClusterOnDieEn	26
6.2.9	MctpEn	26
6.2.10	E2EParityEn.....	26
6.2.11	EarlySnoopEn	26
6.2.12	DegradePrecedence	26
6.2.13	QpiLinkSpeedMode	26
6.2.14	QpiLinkSpeed.....	26
6.2.15	QpiLinkL0pEn.....	26
6.2.16	QpiLinkL1En.....	26
6.2.17	QpiLbEn	26
6.2.18	lIoUniphyDisable	27
6.2.19	QpiLinkCreditReduce	27
6.2.20	QpiConfigTxWci	27
6.2.21	Qpi AdVNACreditArbThreshold	27
6.2.22	QpiCrcMode	27
6.2.23	QpiCpuSktHotPlugEn.....	27
6.2.24	QpiCpuSktHotPlugTopology	27
6.2.25	QpiSkuMismatchCheck	27
6.2.26	PhyLinkPerPortSetting	27
6.2.27	Input Parameter Default Values	27
6.3	OUTPUT PARAMETERS.....	28
6.3.1	CpuInfo.....	28
6.3.2	SysConfig.....	29
6.3.3	WarningLog.....	29
6.3.4	OutLegacyVgaSoc	29
6.3.5	OutIsocEn	29
6.3.6	OutMesegEn	29
6.3.7	OutIsocAzaliaVc1En.....	29
6.3.8	OutClusterOnDieEn.....	29
6.3.9	OutMctpEn	29
6.3.10	OutE2EParityEn	29
6.3.11	OutEarlySnoopEn.....	29
6.3.12	QpiCurrentLinkSpeedMode	29
6.3.13	OutQpiLinkSpeed.....	29
6.3.14	OutQpiLinkL0pEn.....	29
6.3.15	OutQpiLinkL1En.....	29
6.3.16	OutQpiLinkL0rEn.....	29
6.3.17	OutlIoUniphyDisable.....	30
6.3.18	OutQpiCrcMode	30
6.3.19	RasInProgress.....	30
6.3.20	RasEvtType.....	30
6.3.21	RasSocId.....	30



6.3.22	<i>ProgramNonBC</i>	30
6.3.23	<i>ProgramBC</i>	30
6.4	ERROR/STATUS CODES	30
6.5	WARNING LOGS	31
6.6	PORTING GUIDELINES.....	32
6.6.1	<i>QpiHooks.c</i>	32
6.6.2	<i>QpiPlatformEparam.c</i>	32
7	MEMORY REFERENCE CODE.....	33
7.1	BASIC TERMINOLOGIES AND ASSUMPTIONS.....	33
7.2	MRC ENTRY POINT	33
7.3	SETUP OPTIONS AND INPUT PARAMETERS	33
7.3.1	<i>Setup Structure Initialization</i>	34
7.3.2	<i>sysSetup</i>	34
7.3.3	<i>ddrChannelSetup</i>	42
7.4	DIMM POPULATION RULES.....	43
7.5	PLATFORM PORTING AND INTERNAL HOOKS.....	43
7.5.1	<i>MemPlatform.H file</i>	43
7.5.2	<i>MemHooks.c file</i>	43
7.6	POST AND ERROR CODES	43
7.6.1	<i>POST Codes</i>	43
7.6.2	<i>Post Code Stall</i>	4544
7.6.3	<i>Error Codes</i>	45
7.6.4	<i>Warning Log</i>	46
7.7	MRC OUTPUT RESULTS TO PHYSICAL MEMORY	51
7.7.1	<i>Output Structure</i>	51
8	COMPATIBLE DATA STRUCTURE	52
8.1	REQUIREMENTS	52
8.2	PROPOSED SOLUTION	52
8.3	COMPATIBLE ACPI TABLE	52
8.4	SCRATCHPAD	54
8.5	COMPATIBLE BIOS STRUCTURE HEADER.....	54
8.6	BIOS STRUCTURE DEFINITIONS	55
9	FAULTY PARTS TRACKING	59
10	STORAGE FEATURE MRC USERS GUIDE.....	60
10.1	MRC CONFIGURATION OPTIONS	60
10.2	OUTPUT PARAMETERS.....	61
10.3	MEMORY MAP LAYOUT	61



Revision History

Rev. No.	Description	Date
0.50	Initial revision	Sept 2012
0.60	Update Power-on settings	Jan 2013
0.70	Update for V0.70 release	March 12th, 2013
0.72	Update for V0.72 release	April, 2013
0.75	Update for V0.75 release	May 16th, 2013
0.76	Update SysSetup and OptionsExt fields, STS code and major error code	June 26th, 2013
0.77	Revised RasModeEx option.	August 9th, 2013
0.78	Added new option: ATTEMPT_FAST_BOOT_COLD	September 20th, 2013
0.79	Added option fields in OptionsExt (bit11 to bit14). Added a new warning error code 0x25 for " COD Enabled but memory not behind each HA". Added MCODT value selection.	October 4th, 2013
0.80	First Beta release. Added new options and warning errors for dram maintenance test.	October 26th, 2013
0.81	Beta release: added a new option for DLL reset test loop	November 2nd, 2013
0.82	Beta release: update MRC entry point and setup options	December 20th, 2013
0.83	Beta release: update section 10: " Storage Feature MRC Users Guide"	Jan 2014
0.84	Beta release: moved RCOMP_TYPE (bit12) in "option" to RCOMP_TYPE (bit15) in "optionExt" field. Removed "pagetablealias" field in memSetup structure.	February 12 th , 2014
0.90	Production Candidate release. Added new MRC warning code.	March 9 th , 2014
0.96	Production Release. Added one new option: dramaplExtendedRange	March 27 th , 2014

1 Introduction

This document accompanies the Reference Code. The primary purpose of this document is to describe the contents of the Reference Code package. In addition this document describes the integration methods for Independent BIOS Vendors (IBVs) and Intel OEM customers. This document caters to EP processors (HSX-EP and BDX-EP)

1.1 Reference Code package contents

The Reference Code package is a collection of code modules and build tools that can be used to construct an executable BIOS image. It is written primarily in C and ASM programming languages, and it uses the processor cache for stack and data storage.

The code is separated into two categories, Reference code and Sample code. The reference code is used directly in Intel validation and CRB BIOS, and it is expected to be compatible with all platforms that follow the Platform Design Guide. The sample code is used only in the MiniBIOS package and is not expected to be compatible with all platforms. Customers should only modify the sample code.

The reference code package contains:

- Sample code to enter 32-bit protected mode and call into the early reference code module.
- Reference code to initialize the processor(s) for basic operation including microcode loading, MMCFG/bus allocation, cache in No-Eviction Mode (NEM) and other early processor/chipset settings.
- Sample code to find the microcode images, locate the NEM code/data ranges, get early setup options, initialize the serial port, and call into the main reference code module.
- Reference code to initialize the QPI links for proper routing and full speed operation among all processors.
- Reference code to detect DIMMs, initialize each memory controller and establish a system address map.
- Reference code for shared functions like MMCFG access, serial debug messages and other common tasks.
- Sample code to initialize default setup options and platform hooks for use within the reference code.
- Build tools for Portable Executable (PE) files and make files for constructing the MiniBIOS image.

2 MiniBIOS Overview

2.1 MiniBIOS

The reference code package primarily consists of processor initialization, QPI and Memory Reference Code (MRC). The package can be compiled into a binary image called the MiniBIOS for stand-alone execution on the Intel Customer Reference Board (CRB). The image will not be able to boot an OS, but it is capable of executing from the CPU reset vector to the completion of system memory initialization. The MiniBIOS may be used as a sample implementation by IBVs and OEM customers that are developing platforms based on Intel processors.

2.2 Build Tools

The MiniBIOS build requires installation of several tools to compile and link the source code. The minimum tool requirements are listed below:

1. Microsoft ® Macro-Assembler version 6.15 - ml.exe
2. Microsoft ® Segmented Executable Linker version 5.31 - \$(MASMPATH)\link.exe
3. Microsoft ® C to MASM Include File Translator 7.10.3077 - h2inc.exe
4. Microsoft ® 32-bit C/C++ Optimizing Compiler Version 13.10.3077 or later - cl.exe
5. Microsoft ® Incremental Linker Version 7.10.3077 or later - link.exe

The latest versions of these tools are included in the Microsoft ® Visual Studio .NET, with exception of the 16-bit linker and the H2INC utility. The 16-bit linker is included in the MASM installation and must be referenced by the MASMPATH environment variable. The H2INC utility was last included with the .NET 2003 installation.

Functional validation of the reference code will use the latest .NET 2008 versions of the build tools. Support for previous tool versions is limited to build errors only, and does not include execution coverage.

Please note that Microsoft ® Macro-Assembler Version 6.11d exhibits structure packing issues that can break the OEM BIOS interface to MRC. This issue is referenced in Part 7 of the Readme.txt file from the ML6.14.exe package under Microsoft Knowledge Base Article [#228454](#).

2.3 Build options

Memory Build Options:

1. SERIAL_DBG_MSG – When defined enables code for Serial debug messages on Serial port. Default defined
2. LRDIMM_SUPPORT - When defined enables code for LRDIMM support. Default defined
3. INPHI_LRBUF_WA – When defined enables code for LRBUF Workarounds. This will only be enabled if LRDIMM_SUPPORT is enabled. Default defined
4. IDT_LRBUF_WA – When defined enables code for LRBUF Workarounds. This will only be enabled if LRDIMM_SUPPORT is enabled. Default defined
5. BDAT_SUPPORT – When defined enables code for Bios Compatible data structure. Default defined
6. QR_DIMM_SUPPORT - When defined enables code for Quad rank DIMM support. Default defined
7. MARGIN_CHECK – When defined enables code for Dimm Margin checking. Default not defined
8. ME_SUPPORT_FLAG – When defined enables code for ME UMA support. Default not defined
9. XMP_SUPPORT – When defined enables code for Extreme memory profile support. Default not defined.
10. HA2_PLATFORM – When defined enables code for 2 Home Agent\Memory Controller support. Default defined.
11. ASM_INC – When defined enables usage of H2INC utility. Default undefined, need to define for legacy BIOS.
12. LT_STRUCT – When defined enables structures for TPM data. Default undefined.
13. QPI_HW_PLATFORM – When defined enables QPI code for Hardware platform. Default defined.
14. IA32 – When defined enables code to compile reference code in IA32 environment. Default defined. If reference code need to be compiled in X64 then undefine this flag.

2.4 Directory Structure

The following sections of this document assume that the MiniBIOS package is located in the root directory Z:\. While this is not a strict requirement to build the MiniBIOS, it does improve portability of the output files among

multiple hosts. Debug information can be loaded from a remote host without rebuilding source code on a local drive provided that both systems map the root directory to Z:\.

There are two separate methods to create virtual drive Z: depending on physical drive location. If the physical drive is located in the local system, then use the command “SUBST”. If the physical drive is located on another system on the network, then use the command “NET USE”. Examples of each method are given below:

To map virtual drive Z: to a local drive containing the MiniBIOS at C:\PlatformRC\MiniBIOS, use the following commands:

subst Z: /D

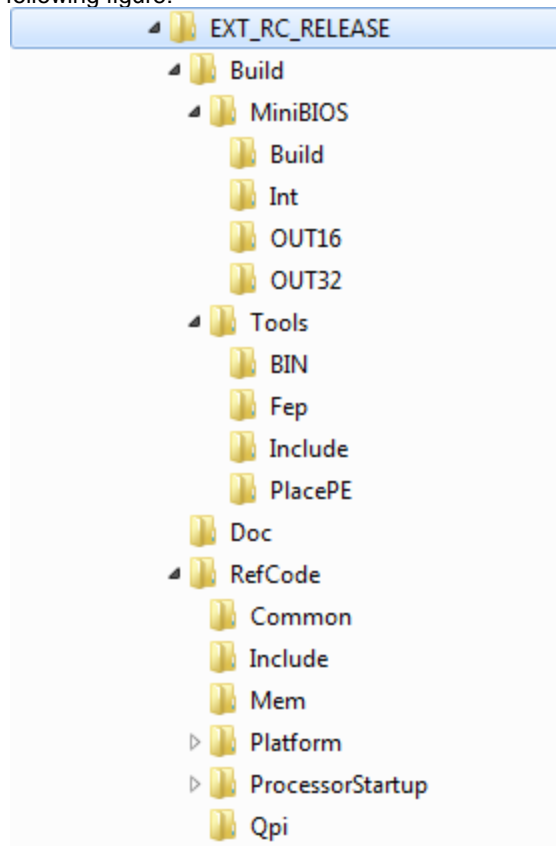
subst Z: C:\PlatformRC\MiniBIOS

To map virtual drive Z: to a network drive containing the MiniBIOS at \\SOXPLO525\c\$\PlatformRC\MiniBIOS, use the following commands:

net use Z: /DELETE

net use Z: \\SOXPLO525\c\$\PlatformRC\MiniBIOS

Given this assumption about the root directory, the subordinate MiniBIOS directory structure is illustrated in the following figure:



A brief description of each directory follows:

1. Build\MiniBIOS: batch and make files to build the MiniBIOS, final binary output files
2. Build\MiniBIOS\OUT16: intermediate 16-bit build output files
3. Build\MiniBIOS\OUT32: intermediate 32-bit build output files
4. Tools\BIN: pre-compiled executables of MiniBIOS tools
5. Tools\Fep: source code for the FindEntryPoint tool
6. Tools\Include: common header files for MiniBIOS tools
7. Tools\PlacePE: source code to the PlacePE tool



8. RefCode: reference code for the high-level flow of QPI and MRC steps, including multi-threaded execution. The following files in this directory are considered “sample code” because they are not used directly in Intel validation BIOS:
 - a. OemMemoryQpiInit.c – platform functions to wrap the reference code
 - b. OemMemoryQpiInit.h – platform header files to wrap the reference code
9. RefCode\Common: reference code to shared functions. This directory should not be modified by customers.
10. RefCode\Include: header files shared by the reference code. This directory should not be modified by customers.
11. RefCode\Mem: reference code to initialize memory. This directory should not be modified by customers.
12. RefCode\Platform\Hooks: platform functions that are called by the reference code. This directory should be modified by the customer as needed.
13. RefCode\Platform\Include: platform header files that are used by reference code. This directory should be modified by the customer as needed.
14. RefCode\Platform\Setup: platform default input parameters to the reference code. This directory should be modified by the customer as needed.
15. RefCode\ProcessorStartup: reference code for early processor initialization and basic chipset settings that are prerequisite for QPI and MRC. The following files in this directory are considered “sample code” because they are not used directly in Intel validation BIOS:
 - a. CpuPm32.asm – Switch to 32-bit protected mode and call ProcessorStartup module
 - b. CpuMain.c – Serial port init, reserve stack frame and call reference code module
 - c. CpuStart.asm – Call CpuMain, check stack usage, tear down NEM and stall
 - d. ProcessorStartupOem.asm – platform hooks for early init
 - e. ProcessorStartupPlatform.inc – platform build options for early init
16. RefCode\Cpu\Ucode: microcode patches loaded by ProcessorStartup module.
17. RefCode\Qpi: reference code for QPI initialization among all processors. This directory should not be modified by customers.
18. Doc*: this file and readme.txt files that track changes to the reference code.

2.5 Build Environment

This section describes the environment variables that are required to build the MiniBIOS.

1. Open a .NET DOS command box.
2. Run the VSvars32.BAT file that is supplied with Microsoft Visual Studio .NET. This will modify the search path to run the compiler and linker.
3. “SET MASMPATH = <Path to Microsoft Segmented Executable Linker>”. This is required to link the object components into a binary executable.
4. Optionally “SET BIOS86DIR = <Path to the Microsoft Visual Studio .NET tools>\”. This option is useful if the .NET tools are not included on the search path.
5. Optionally “SET MINITOOLS = <Path to the H2INC utility>\”. This option is useful if the H2INC utility is not included on the search path.

2.6 Building the MiniBIOS

Change the working directory to: Z:\Build\MiniBIOS

To do a full build of the MiniBIOS image, execute the commands:



BuildMiniBIOS clean

BuildMiniBIOS

To do a partial build of the MiniBIOS, execute the commands:

BuildMiniBIOS

This generates the 1 MB MiniBIOS binary file Z:\Build\MiniBIOS\BIOS.BIN.

2.7 Debugging the MiniBIOS Using Intel SoftSimulation

This section provides a few quick tips for debugging the MiniBIOS image using the SoftSimulation environment. Please refer to the SoftSimulation User's Guide for more detailed instructions.

2.7.1 Configuring SoftSimulation to Load the BIOS image file

The name of the MiniBIOS binary file to be updated in simics file which is to be ran. Simics files are located at C:\Users\ldSid\simics-workspace\targets\Grantley. This file also will be used to configure system configuration information, including number of DIMMs, processor sockets, cores, threads, etc.

Please refer to simics release note for target script configuration.

In this example, the name of the MiniBIOS binary file is Z:\Build\MiniBIOS\BIOS.BIN.

2.7.2 Symbolic Debugging of MiniBIOS

Start the simulator by running the Simics executable. After the simulator completes loading, an ITP window will be visible to accept command input. Open the source code window to make sure that ITP is halted at the CPU reset vector. In the ITP command window, type the following commands:

```
go til 0xffff00640p; wait
load Z:\Build\MiniBIOS\out32\procstartupinit.bin nocode
load Z:\Build\MiniBIOS\out32\rcinit.bin nocode
```

The command "go til 0xffff00640p" instructs the simulator to execute until the physical address 0xffff00640 is reached. This address is the entry point, ProcessorStartup, of the PROCSTARTUPINIT.BIN module. It can be found by building the MiniBIOS and reading the file Z:\Build\MiniBIOS\out32\procstartupentry.inc. This entry point is also displayed on the screen at the end of the MiniBIOS build. It can change from build to build, so make sure you use the correct address in the command line.

The command "load Z:\Build\MiniBIOS\out32\procstartupinit.bin nocode" loads the symbolic debug information into the ITP source-level debug environment. After this command has executed, C and Assembler source code, as originally written by the BIOS author, may be viewed in the source code window and stepped through.

The command "load Z:\Build\MiniBIOS\out32\rcinit.bin nocode" loads the symbolic debug information into the ITP source-level debug environment. The other key entry point, MemoryQpInit in the RCINIT.BIN module, can be found by reading the definition for RCINITENTRYPOINT in the file Z:\Build\MiniBIOS\out32\rcentry.h.

2.7.3 Using the DUMPBIN utility

DUMPBIN.EXE is a utility that comes with Microsoft Visual Studio .NET to dump information from PE/COFF binary file formats. A sample of the output is shown below:

```
Z:\Build\MiniBIOS\OUT32>dumpbin rcinit.bin /headers
Microsoft (R) COFF/PE Dumper Version 7.10.3077
Copyright (C) Microsoft Corporation. All rights reserved.
```

Dump of file rcinit.bin

PE signature found

File Type: DLL

FILE HEADER VALUES

```
14C machine (x86)
3 number of sections
0 time date stamp Wed Dec 31 19:00:00 1969
0 file pointer to symbol table
0 number of symbols
E0 size of optional header
2102 characteristics
```



Executable
32 bit word machine
DLL

OPTIONAL HEADER VALUES

10B magic # (PE32)
9.00 linker version
34340 size of code
8F40 size of initialized data
0 size of uninitialized data
336C0 entry point (FFF536C0)
240 base of code
34580 base of data
FFF20000 image base (FFF20000 to FFF5D4BF)
20 section alignment
20 file alignment
5.00 operating system version
0.00 image version
5.00 subsystem version
0 Win32 version
3D4C0 size of image
240 size of headers
0 checksum
3 subsystem (Windows CUI)
400 DLL characteristics
No safe exception handler
100000 size of stack reserve
1000 size of stack commit
100000 size of heap reserve
1000 size of heap commit
0 loader flags
10 number of directories
0 [0] RVA [size] of Export Directory
0 [0] RVA [size] of Import Directory
0 [0] RVA [size] of Resource Directory
0 [0] RVA [size] of Exception Directory
0 [0] RVA [size] of Certificates Directory
3C7A0 [C60] RVA [size] of Base Relocation Directory
36810 [1C] RVA [size] of Debug Directory
0 [0] RVA [size] of Architecture Directory
0 [0] RVA [size] of Global Pointer Directory
0 [0] RVA [size] of Thread Storage Directory
0 [0] RVA [size] of Load Configuration Directory
0 [0] RVA [size] of Bound Import Directory
0 [0] RVA [size] of Import Address Table Directory
0 [0] RVA [size] of Delay Import Directory
0 [0] RVA [size] of COM Descriptor Directory
0 [0] RVA [size] of Reserved Directory

SECTION HEADER #1

.text name
34321 virtual size
240 virtual address (FFF20240 to FFF54560)
34340 size of raw data
240 file pointer to raw data (00000240 to 0003457F)
0 file pointer to relocation table
0 file pointer to line numbers
0 number of relocations
0 number of line numbers
60000020 flags
Code
Execute Read

SECTION HEADER #2

.data name
8217 virtual size
34580 virtual address (FFF54580 to FFF5C796)
8220 size of raw data
34580 file pointer to raw data (00034580 to 0003C79F)
0 file pointer to relocation table



0 file pointer to line numbers
0 number of relocations
0 number of line numbers
C0000040 flags
Initialized Data
Read Write

Debug Directories

Time	Type	Size	RVA	Pointer
4B845ABA	cv	3B	0003C75C	3C75C
Format: RSDS, {75E96270-0E9B-4592-9348-97AA4EDC851D}, 1, Z:\Build\MiniBIOS\OUT32\RCINIT.pdb				

SECTION HEADER #3

.reloc name
D18 virtual size
3C7A0 virtual address (FFF5C7A0 to FFF5D4B7)
D20 size of raw data
3C7A0 file pointer to raw data (0003C7A0 to 0003D4BF)
0 file pointer to relocation table
0 file pointer to line numbers
0 number of relocations
0 number of line numbers
42000040 flags
Initialized Data
Discardable
Read Only

Summary

8220 .data
D20 .reloc
34340 .text

The entry point of the PE/COFF executable is indicated in the line:

336C0 entry point (FFF536C0)

This address can be used in itp_ssdrv.cfg to automatically halt and load source level debug information. The above dump is only a sample and the actual address may differ.

3 Details of MiniBIOS

This section describes the MiniBIOS layout in detail. BIOS engineers should read this section carefully, as it explains how to build the reference code modules and integrate them into a customer BIOS.

3.1 MiniBIOS layout

The layout of the MiniBIOS is shown in the figure below. The sample CPUPM32.BIN module occupies the address range based at (4GB – 64KB). This area is what would be the customer's main BIOS. The reference code occupies the address range based at (4GB – 2MB), extending up to the size of the PE executable modules PROCSTARTUPINIT.BIN and RCINIT.BIN. The final output image BIOS.BIN is 2MB in size and top-aligned to (4GB – 2MB).

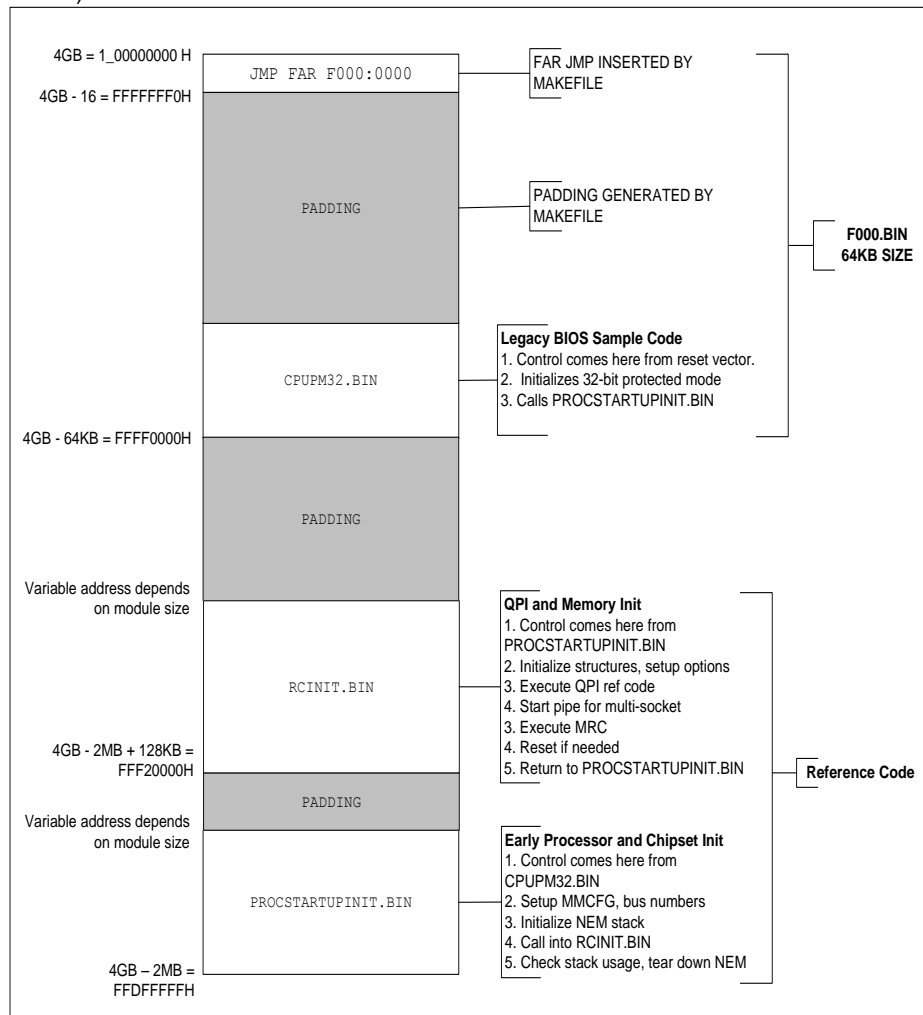


Figure 1: MiniBIOS Layout

3.2 Building the MiniBIOS

The MiniBIOS is built by executing the BuildMiniBIOS.bat file within the Z:\Build\MiniBIOS directory, which in turn launches nmake on the local makefile. The build occurs in the following sequence:

1. Compile and link reference code into the OUT32\RCINIT.DLL module. The entry point is *MemoryQpiInit*.



2. Run **PlacePE** to locate the RCINIT.BIN module at the 32-bit physical address RC_ROM_BASEADDRESS specified in BuildOptions.mak. The address must be aligned to a 16-byte boundary.
3. Run **FindEntryPoint** to get the entry point of the RCINIT.BIN module and save rcentry.h
4. Compile and link processor startup reference code into the OUT32\PROCSTARTUPINIT.DLL module. The entry point is *ProcessorStartup*.
5. Run **PlacePE** to locate the PROCSTARTUPINIT.BIN module at the 32-bit physical address PROCSTARTUP_ROM_BASEADDRESS specified in BuildOptions.mak. The address must be aligned to a 16-byte boundary.
6. Run **FindEntryPoint** to get the entry point of the PROCSTARTUPINIT.BIN module and save procstartupentry.inc.
7. Compile and link the sample code into the CPUPM32.BIN module. This would normally be linked with other OEM BIOS components to form the F000 ROM image.
8. Concatenate the binary files with padding to create the BIOS.BIN image. This file is the complete 1MB image that is shown in *Section 3.1 Mini-BIOS Layout*.

3.3 Execution Flow

This section describes the execution flow of the MiniBIOS as a sample implementation.

1. Control comes from the reset vector at 0xFFFFFFF0 to *_PM32Entry* in RefCode\ProcessorStartup\CpuPm32.asm.
2. Switch to 32-bit protected mode and call into the PROCSTARTUPINIT module. The entry point is *ProcessorStartup* in RefCode\ProcessorStartup\ProcessorStartup.asm.
3. Initialize MMCFG, bus numbers, early chipset settings, NEM stack and call into the RCINIT.BIN module. The entry point is *MemoryQpInit* in RefCode\MemoryQpInit.c
4. Initialize QPI input parameters, execute QPI reference code, and process output parameters.
5. Setup the pipe environment for multi-socket execution, initialize MRC input parameters, execute MRC, and process output parameters. A system reset will be issued if necessary.
6. After MRC completes, copy reference code output parameters to system memory. Disable the pipe handler, tear down NEM stack on the remote sockets and return to the PROCSTARTUPINIT module.
7. Check the stack usage, tear down NEM stack on the SBSP and stall execution with an infinite loop (JMP \$).

3.4 MiniBIOS Tools

This section describes the build tools that come with the reference code.

3.4.1 FindEntryPoint

This utility parses the header of the PE binary file and calculates the 32-bit physical address corresponding to the entry point into the binary. It accepts one optional parameter to specify the output language, /C for C code or /A for ASM code.

Syntax: FindEntryPoint / [C | A] infile

infile: PE file whose entry point is to be found.

Example:

findentrypoint /C out32\rcinit.bin

Output:

#define RCINITENTRYPOINT 0xffff536c0

3.4.2 PlacePE

This utility locates a PE binary file at a user-supplied 32-bit physical address where it will execute from. It has the ability to locate the executable at any byte-boundary inside 4GB address space, however the executable should normally be aligned to a 16-byte address boundary.

Syntax: placepe [infile] [outfile] [base address]

infile: PE file that is to be fixed up.



outfile: The fixed-up PE file that is generated.

base address: The 32-bit flash address where this code resides.

Example:

```
copy out32\rcinit.dll out32\rcinit.bin
```

```
Placepe out32\rcinit.bin out32\rcinit.bin 0xffff20000
```

Note: This utility requires that the infile be copied to the outfile prior to execution. The outfile will then be patched with appropriate data by this utility.

4 Common Reference Code Components

4.1 Overview

This section details information that is shared or referenced by multiple components.

4.2 BIOS Reference Code High Level Execution Flow

This section shows a high level initial cold-boot flow of the BIOS reference code components. Details are present in the following chapters

- CPU Initialization BIOS Reference Code / Early Initialization section.
- QPI Reference Code ("Phase 1": initializes registers that require a reset to latch, warm reset then runs "phase 2")
- Pipe Init
- Memory Reference Code ("Phase 1": typically exits with a frequency change on reset request).
- CPU Initialization BIOS Reference Code / Feature Early Configuration (sets registers that require a reset to take effect)
- Pipe Exit
- Check for Reset (Typically a reset is required after the first pass through the early BIOS), and issue reset if needed.

Below shows a high level BIOS reference code flow between the BIOS reference code components after software generated reset. Details are present in the following chapters

- CPU Initialization BIOS Reference Code / Early Initialization section.
- QPI Reference Code ("Phase 2")
- Pipe Init
- Memory Reference Code ("Phase 2": system memory initialization).
- CPU Initialization BIOS Reference Code / Feature Early Configuration
- Pipe Exit
- Check for Reset (not typically needed at this point in the flow)

4.3 Assumptions

The various BIOS reference code components assume that the system is initialized at roughly the same level as the miniBIOS and the components are called in the same order. For example, for the Memory Reference code to run multi-socket, it requires the "pipe" infrastructure to be initialized, and for the pipe to work, it requires the QPI links to be initialized at least to the point of Phase 1 of the QPI reference code, and phase 1 of the QPI reference code requires the CPU bus numbers and configuration space be initialized as in the early initialization section of the CPU BIOS reference code.

4.4 Platform Porting and Reference code Hooks

Other files in the Refcode\Platform subdirectories must also be ported for a given platform. There are also files in the RefCode directory and subdirectories should not be modified. By default, these files implement settings compatible with the Intel customer reference board and the Intel Platform Design Guide.

4.4.1 Options

Common Option fields.

4.4.1.1 PROMOTE_WARN_EN (BIT0)

When set, all warnings logged are treated as fatal errors.

4.4.1.2 PROMOTE_MRC_WARN_EN (BIT1)

When set, Memory Reference Code warnings are treated as errors independently of PROMOTE_WARN_EN.



4.4.1.3 HALT_ON_ERROR_EN (BIT2)

When set, all errors stop at a halt loop to allow for easier debug when an error occurs.

4.4.2 lowGap

This is a byte size field that contains the desired size of the low MMIO gap in units of 64MB. Low MMIO gap is defined by TOLM to 4 GB. Minimum size is 512MB assuming PCI configuration space MMCFG is below 4GB. This is a system address region right below 4GB typically used for MMIO configuration registers, memory resources allocated to PCI devices, and so on.

4.4.3 highGap

This is a byte size field that contains the desired size of the high MMIO gap in units of 64MB. Minimum size is 64 MB. HighGap is defined by (LIMIT + 10h - TOLM) to 64 GB. This option is included for systems with special needs and require such memory gap or hole be reserved above 4GB.

4.4.4 serialDebugMsgLvl

Specifies what level of debug messages will be sent to serial port. Available options are a bitfield where bit 0 set = minimal information, bit 1 set = more information (but longer execution time), and bit 3 set = a full trace of the BIOS configuration register accesses.

```
#define SDBG_MIN          BIT0
#define SDBG_MAX          BIT1
#define SDBG_TRACE       BIT2
#define SDBG_MEM_TRAIN   BIT3 + SDBG_MAX
#define SDBG_TST         BIT4
#define SDBG_MINMAX      SDBG_MIN + SDBG_MAX
```

4.4.5 bsdBreakpoint

Specifies the BIOS serial debug breakpoint. It is used to stall BIOS execution at a predefined checkpoint so that EV tools can be run in the proper context (before switching to normal mode and locking registers). Format is similar to port 80 postcode.

4.4.6 maxAddrMem

Maximum addressable memory supported by the platform. Haswell Processor supports up to 46-bit addressing. This input should be the total number of addressable bytes in 256MB units. (0x40000 for 46-bit and 0x1000 for 40-bit). 46-bit addressing is the default.

4.4.7 SocketType

Type of socket used. Socket R3 (= 0) for Platform.

4.4.8 debugPort

User configurable IO port for post code which is traditionally located at 0x80. 0x80 will be used if this field is 0.

4.4.9 nvramPtr

32-bit pointer to an optional OEM NVRAM image to be copied into the host NVRAM structure.

4.4.10 sysHostBufferPtr

32-bit pointer to an optional OEM provided Host structure.

4.4.11 mmCfgBase

Base address of low PCIEX MMIO in 256MB granularity.

4.4.12 mmCfgSize

Size of low PCIEX MMIO memory range. Must be 64MB (0x4000000), 128MB, or 256MB (default).

4.4.13 mmioIBase

Low MMIO base, 64M aligned.

4.4.14 mmioISize

Low MMIO size, 64M aligned.

4.4.15 mmiohBase

The hex value of the address lines [45:32] where MMIO High resources should be placed.

4.4.16 mmiohSize

Number of 1GB contiguous regions to be assigned for MMIOH space per CPU.

4.4.17 pchUmaEn

UMA support for PCH.

4.4.18 numaEn

Set to 1 to enable NUMA, or 0 to disable (UMA).

4.4.19 isocEn

Set 1 to enable Isoc

4.4.20 DciVcm

Set 1 to enable DciVcm, or 0 to disable.

4.4.21 MeSegEn

Set 1 for MeSeg Support. Exclusive with IsochEn

4.4.22 dcaEn

Set to 1 to enable DCA, or 0 to disable.

4.4.23 PSMIEnabledSupport

Set 1 to enable PSMI.

4.4.24 BdatEn

Enable BDAT information for OS consumption

4.4.25 consoleComPort

Base address of the console ComPort if it's available in the system.

4.4.26 debugComPort

Base address of debug ComPort if it's available in the system.

4.4.27 usbDebugPortSetup

A structure with information about the platform USB Debug compliant port, if one is available in the system.



4.5 Scratch pad register usage

Haswell contains 8 sticky, 7 non-sticky lock bypass and 16 non-sticky scratchpad registers for BIOS usage. The reference code takes advantage of several of these registers which are documented here. Some of these registers have been set aside for future development and have been marked as “Intel Reserved”. All other registers are marked as available customer usage. Please refer to Scratch-Pad consumption documentation in Doc\ScratchPadList.xls

4.6 High and Low MMIO BAR allocation

QPIRC assigns High MMIO (above 4G) and Low MMIO (below 4G) for each socket. QPIRC chooses 0x3800_000_0000 as default address for High MMIO and 0x8000_0000 as default address for low MMIO. This start address can be changed by OEM to meet its platform needs.

The highest address supported by Haswell is 0x4000_0000_0000 (2⁴⁶). Address range 3800_0000_0000 to 4000_0000_0000 is available for OEM to assign for high MMIO space. High MMIO Base address for non-BSP will change based on size selected. To specify the size, OEM could simply change *mmiohSize* in *commonSetup* structure with increment of 1GB to any size limited to the number of sockets within the range between 3800_0000_0000 to 4000_0000_0000. For example, in QPIRC OEM can specify up to 1TB of high MMIO space for each socket in 8 socket glue-less configuration. Similarly OEM can specify up to 2TB of high MMIO space for each socket in 4 socket glue-less configuration.

Code snippet in reference code to change the size in the reference code is as below:

```
Filename: CommonSetup.c
Routine: GetSetupOptions
setup->common.mmiohSize = 2; // 2G per IIO
```

Code snippet in QPIRC to change start address of High MMIO base is as below:

```
File: QpiLib.c
Routine: AllocateResources
//
// Initialize the Base value of the resources
//
BusBase = 0;
IoBase = 0;
IoApicBase = IOAPIC_BASE;
MmioLBase = host->setup.common.mmCfgBase + host->setup.common.mmCfgSize;
//100_0000_0000 = 1TB per socket for 8SG
MmiohBase.lo = 0;
MmiohBase.hi = 0x3800; // MMIOH_BASE;
```

After QPIRC allocates and divides MMIO resources across all sockets standard PCI enumeration code in platform codebase will assign these resources based on PCI configuration. Each end point device will request these resources through standard PCI BAR registers. In case platform BIOS uses EDKII core code PCI enumeration and MMIO resource allocation could be found in File “PciEnumatorSupport.c” Function name “PciParseBar”

5 CPU Initialization Reference Code

5.1 Overview

CPU init Reference Code contains source files in Intel x86 assembly and C languages. It exemplifies BIOS boot flow from the reset vector to the end of MRC, when system memory has been initialized and becomes available. It depends on QPI RC for CPU population topology discovery and employs a cold-boot/warm-boot 2-pass approach to complete the basic CPU, QPI and memory initialization.

5.2 CPU Reference Code Execution Flow

This section shows a high level flow of the CPU reference code and its relation to the QPI/Memory reference code as well as what is being done by the CPU reference code where applicable.

Early CPU Init

Run by all Package BSPs

- Switch to 32-bit protected mode
- Load microcode update
- Configure DCU_Mode per user input
- Enable No Eviction Mode (NEM)
- Config TARGET_LIST, CPUBUSNO and MMCFG for CSR access
- Select SBSP (legacy socket)
- Non-SBSPs check-in and loop on Scratchpad CSR

Invoke QPI RC to discover CPU population topology

- On cold-boot path, CPU population and final CPUBUSNO, MMCFG information is determined and stored in sticky SR, and reset request is raised.

PIPE Init

- Set up a mailbox mechanism by which SBSP (PIPE Master) can send command/data to other BSPs. This is intended for multi-threaded MRC execution.

CPU Feature Early Configuration

- Config FlexRatio, DesiredCores, SMT, etc that need a reset to take effect
- Run by SBSP only, via CSR access to all CPUs present
- Set CPU reset global flag if needed

Invoke MRC to initialize memory

PIPE Exit

- SBSP (PIPE Master) sends command via PIPE mailbox to other BSPs to exit PIPE. SBSP sends INIT-IPI to put all BSPs in Wait-For-SIPI state.

Check for Reset

- Consolidate/minimize potential multiple resets
- Check global flag for reset request by all sources
 - QPIRC
 - MRC
 - CPU Feature Early config
- Initiate the type of reset requested in the priority order below:
 - Power Good Reset (IO Port 0xCF9 = 0xE)
 - Warm Reset (IO Port 0xCF9 = 0x6)

5.3 Assumptions

The following assumptions are used by the CPU Ref Code:

- Max of 4 CPU sockets is validated on Intel platforms.
- Reference code will support up to 8 sockets.
- Reference Code has a limitation that MAX_CPU_SOCKETS must be either 4 or 8.
- 8 sockets support is build time decision.
- Legacy socket (with PCH connection via DMI link) is always present and designated as system BSP.
- Reserved usage of mmx0-7, xmm1-7 CPU registers in SEC phase
- Entire reference code needs to be cached into NEM code space. Early reset code before enabling NEM space is exception to this condition.

5.4 Platform BIOS Integration Guide

CPU init Reference Code provides platform/OEM customization provisions such as Build-time defines/equates, input parameters, and OEM code hooks. This chapter contains a summary of platform/OEM customization files and hooks related to CPU init reference code.

5.4.1 Build-Time Equates and Defines

Filename: **ProcessorStartupPlatform.inc**

The following equates define PCIe config base address and No Eviction Mode parameters:

```
PCIEX_BASE_ADDRESS EQU 08000000h
; MMCFG / PCIe Config Base Address, used as cold boot default only.
; Use EDKII_GLUE_PciExpressBaseAddress in PlatformHost.h file to
; define platform choice of PCIe config base address
```



```

CODE_REGION_BASE_ADDRESS      EQU 0ffe00000h
CODE_REGION_SIZE              EQU 0200000h
CODE_REGION_SIZE_MASK        EQU (NOT (CODE_REGION_SIZE - 1))
DATA_STACK_BASE_ADDRESS      EQU 0fe080000h
DATA_STACK_SIZE              EQU 080000h
DATA_STACK_SIZE_MASK        EQU (NOT (DATA_STACK_SIZE - 1))
MAX_CPU_SOCKETS              EQU      x
    ; where default x = 4 for 4S support
    ; to build 8S, change to 8

```

5.4.2 Input Parameters

CPU RC related input parameters are defined in file CpuHost.h, struct cpuSetup, as shown below:

```

struct cpuSetup {
    // CPU init related user options
    UINT8  dcuModeSelect;    // 0: 32KB 8-way (hardware default). Non-zero: 16KB 4-way with ECC (CPU MSR 031h)
    UINT8  flexRatioEn;     // FLEX_RATIO Override Enable. If non-zero then flexRatioNext will be used
    UINT8  flexRatioNext;   // Target FLEX_RATIO, common for all CPU sockets.
    UINT32  CoreDisableMask[MAX_SOCKET]; // CoreOffMask value for each CPU socket. If 0 then no CPU cores
                                           // will be disabled by BIOS
    UINT8  smtDisable;      // 0/1 Disable/Enable SMT(HT). common for all CPU sockets
    UINT8  vtEnable;       // 0/1 Disable/Enable VMX. Common for all CPU sockets
};

```

Other common input parameters that affect CPU RC are defined in SysHost.h, struct commonSetup

```

struct commonSetup {
    UINT32  sysHostBufferPtr; // address to copy sysHost structure to after MRC.
    UINT32  mmCfgBase;       // MMCFG Base address, must be 64MB aligned
    UINT32  mmCfgSize;       // MMCFG Size address, must be 64M, 128M or 256M
    UINT8   dcaEn;           // 1 - Enable; 0 - Disable
};

```

5.5 Routines for Platform Customization

Filename: **ProcessorStartupOem.asm**

Routine name: **oem_get_cpu_rom_setup_options PROC**

Returns: Build-time platform options related to CPU early init

- Called early after reset, before loading ucode
- Returns ucode region address/size, etc
- Sample implementation assumes ucode update files are included in the source code as .txt files (similar to Nhm reference code)

Routine name: **OemFindMicrocode PROC**

Returns: 32-bit linear pointer to the CPU ucode patch header

- Called early after reset, before loading ucode
- Sample implementation similar to Nhm reference code.

Routine name: **GetPlatformCpuSettings PROC**

Returns: Run-time platform options related to CPU early init

- Called early after reset, before NEM init
- Returns DCU_MODE select, which is usually determined by the input parameter setup->cpu.dcuModeSelect.
- Sample implementation based on intel ref code package implementation

Filename: **CommonSetup.c**

Function Name: **GetSetupOptions()**

- Provide common default setup options.
- Called after early power-on NEM init, before QPI RC
- mmcfgBase. This is the value for setting the final MMCFG(PCIe) config space base address on the warm boot path.
- maxBusNum can be 0xFF, 0x7F, or 0x3F

Filename: **OemMemoryQpiInit.c**



Function Name: **OemInitializePlatformData()**

- Provide OEM platform specific setup options and/or to override default input options.
- Called after early power-on NEM init, before QPI RC
- Examples include pointer to memory buffer where the sysHost structure will be copied to after memory has been initialized, CPU Flex Ratio, Desired_Cores, SMT enable/disable, DCU_Mode, etc.

Function Name: **OemMemoryInit()**

- OEM hook called after MRC execution
- Sample implementation copies the sysHost structure to a specified buffer in system memory (default address = 0x00100000).

5.6 POST Codes and Serial Port Debug Messages

5.6.1 POST Code at IO Port 0x80

Below is the POST code output to IO port 0x80.

Post Code (Hex)	Status	Error
01	Uncore BUS/MCFG configuration	
02	Local Microcode	
CD		Microcode Not found
03	Enable NEM mode	
D0		NEM Data Stack Test Failed

5.6.2 Serial Port Debug Messages

Below is the capture of a typical example of debug messages to serial port.

Emulation Value is: 1!

Running on SoftSDV

Revision: 32

QPI Init starting...

(QPI RC output messages)

....

....

completed!

Pipe Init starting...completed!

Memory Init starting...completed!

CPU Feature Early Config starting...completed!

Pipe Exit starting...completed!

Emulation Value is: 1!

Running on SoftSDV

Revision: 0

QPI Init starting...

(QPI RC output messages)

....

....

completed!

Pipe Init starting...completed!

Memory Init starting...

(Memory RC output messages)

....

.....

MRC is done!

sysHost structure @ FEFED18 has been copied to memory address 00100000, size = 000102A0

completed!



```
CPU Feature Early Config starting...completed!  
Pipe Exit starting...completed!  
Checking for Reset Requests ... None Continue with system BIOS POST...
```


6 QPI Reference Code

6.1 Overview

This Platform allows systems with variety of QPI topologies. The QPIRC code dynamically determines all aspects of the topology such as number of and type of sockets populated and the link that connects them together. However system aspects such as link electrical parameters are system specific. Also parameters such as Low MMIO region base and size are configurable for these platforms. So QPIRC provides hooks that can be modified to return the data that are system specific. It also provides hooks related to Error/Status code display, issuing reset, changing the input parameters and so on. This section provides the details of input parameters that can be changed by platform BIOS, output parameters produced by QPIRC and the porting needed to integrate QPIRC into platform BIOS code base.

6.2 Input Parameters

The QPIRC uses a set of input parameters to configure how it initializes the QPI fabric of the system. It initializes the input parameters with known good default values. It then calls an OEM hook to override those default values with platform specific values. Typically these values are stored in NVRAM and they are changeable thru BIOS setup menus. The format of the input parameters structure may vary with each QPIRC release. The individual BIOS developer must review this structure to ensure that all fields are properly initialized. The individual field definitions for QPIRC are given below. Some of the input parameters are not applicable to EP platforms but to enable code sharing with EX platforms, they retained in EP QPIRC. Please refer to QpiHost.h file for more info about these fields.

6.2.1 PPINrOptIn

Unused.

6.2.2 BusRatio/IoRatio/MmioIRatio

These parameters are specified by three byte arrays of [MAX_CPU_SOCKETS] to indicate the ratio of Bus/IO/MMIOL resources to be allocated for each socket. Any nonzero value for these parameters indicates the QPIRC to allocate resources for that socket. If resource is requested for a socket that is not currently populated, QPIRC will assume that the ratio is 0 for that socket and won't allocate any resources for it. If resource is not requested for a socket that is populated, QPIRC will force the ratio for that socket to 1. Maximum value allowed is 8 for Bus/Io resources; for MmioI, the maximum value allowed is (4G – MMIOLBase) / 64MB. Note that the size of BUS/MMIOL range is determined based on other input options provided to QPIRC. Size of Legacy IO is fixed at 64KB. For a given socket, multiplying the resource ratio with the chunk size provides the total resource size allocated to that socket. The resource for Bus, IO & MMIOL ranges is 32MB (32 buses), 8KB & 64MB, respectively.

For example, in a 4S system, user can select the IO ratio as 5:1:1:1 in which case, the 64KB IO region will be assigned as follows:

```
Socket 0: 5 * 8K = 40K (0x0000 - 0x9FFF)
Socket 1: 1 * 8K = 8K (0xA000 - 0xBFFF)
Socket 2: 1 * 8K = 8K (0xC000 - 0xDFFF)
Socket 3: 1 * 8K = 8K (0xE000 - 0xFFFF)
```

If equal distribution is not possible, socket with lowest socket ID will get additional chunks. For example, if the ratio is 4:1:1:1, then the resource allocation will be as follows:

```
Socket 0: 40K (0x0000 - 9FFF) ((4+1)*8K)
Socket 1: 8K (0xA000 - 0xBFFF)
Socket 2: 8K (0xC000 - 0xDFFF)
Socket 3: 8K (0xE000 - 0xFFFF)
```

6.2.3 LegacyVgaSoc

This option conveys which socket should be set as the target for legacy video range. Default is to set the legacy socket as the target for this range. This is based on socket ID. If the socket ID selected by this parameter is not actually populated, then legacy socket will be forced as Legacy VGA target.

6.2.4 MmioP2pDis

This option controls the P2P traffic across sockets. This doesn't affect P2P traffic whose source and target end points belong to the same IIO bus hierarchy. Value 0 will enable the P2P; Value 1 will disable the P2P. Other values are invalid.



6.2.5 IsocAzaliaVc1En

This option controls, when Azalia is enabled, if it uses VCp. Value 1 will use VCp for Azalia traffic; Value 0 will use VC1 for Azalia traffic; Value 2 will be treated as Auto in which case VCp will be used for all topologies that support ISOCH transactions.

6.2.6 DebugPrintLevel

Bit map identifying what level of debug message output is needed.

- Bit 0 - Fatal
- Bit 1 - Warning
- Bit 2 - Info Summary
- Bit 3 - Info detailed.

Value of 1 in a bit enables corresponding debug message level; Value of 0 disables it.

6.2.7 AltRtid2S

Unused

6.2.8 ClusterOnDieEn

This option enables the Cluster on Die feature. 1 to enable, 0 to disable. This only is only available in 1/2S EP.

6.2.9 MctpEn

Enable MCTP feature. 1 to enable, 0 to disable. Default is disable.

6.2.10 E2EParityEn

This option enables end to end parity. 1 to enable, 0 to disable. Default is disable.

6.2.11 EarlySnoopEn

This option enables early snoop. 1 to enable, 0 to disable, default is disable.

6.2.12 DegradePrecedence

This option uses DEGRADE_PRECEDENCE definition. TOPOLOGY_PRECEDENCE is default

6.2.13 QpiLinkSpeedMode

This option selects Slow vs Full Speed mode for all QPI links in the system. A value of 0 indicates to remain in slow mode. A value of 1 selects to transition to full speed mode.

6.2.14 QpiLinkSpeed

This option selects the desired speed of all QPI links in the system in Full Speed mode. This option is system wide. A value of 1 selects 6.4GT/s. A value of 3 selects 8.0GT/s, and a value of 5 selects 9.6GT/s. To auto select the fastest mode, input a value of 6. QpiHost.h provides equates that can be used such as SPEED_REC_80GT and MAX_QPI_LINK_SPEED.

6.2.15 QpiLinkL0pEn

To enable L0p feature which cuts the link width in half when bandwidth utilization drops below a configurable low threshold. Default value is 0 (disable).

6.2.16 QpiLinkL1En

To enable L1 feature which is similar to L0s, but also shuts down the clock lane, saves more power and takes longer to restart. Default value is 1 (enable).

6.2.17 QpiLbEn

To enable QPI link loopback mode. Value 0 is disable (default), 1 is enable.

6.2.18 lioUniphyDisable

Byte array of [MAX_CPU_SOCKETS]. This option is valid only for non-legacy socket. It is used to clock gate PCIe ports and IIO logic on platforms that don't have the slots implemented on any of the CPU sockets. If the platform doesn't have PCIe ports implemented but does use the IIO VPP for memory hotplug then it should be set to 2; otherwise it should be set to 1.

6.2.19 QpiLinkCreditReduce

To reduced the number of link layer VNA credits for system and per link. Should be used only for debug.

6.2.20 QpiConfigTxWci

Unused

6.2.21 Qpi AdVNACreditArbThreshold

To program R3QPI credit merging configuration field.

6.2.22 QpiCrcMode

Unused.

6.2.23 QpiCpuSktHotPlugEn

Indicates if Skt is hot plug enabled. 0 is disable (default), 1 is enable.

6.2.24 QpiCpuSktHotPlugTopology

0 is 4S topology, 1 is for 8S topology

6.2.25 QpiSkuMismatchCheck

Option to check Sku mismatch. 0 is No, 1 is yes (default)

6.2.26 PhyLinkPerPortSetting

Data structure of [MAX_CPU_SOCKETS]. Except QpiLinkSpeed, all members of the structure provide input parameters per link for debug purposes. QpiLinkSpeed, is used to configure the per link speed setting. Note that link 0 & 1 must be operating at the same speed while link 2 can have different speed. The following fields are part of the structure:

- **QpiPortDisable:** To disable or enable this port.
- **QpiLinkCreditReduce :** To force reduced link credit operation, set a value of 1 in the QpiLinkCreditReduce field.
- **QpiProbeType :** To specify a midbus type probe, set a value of 1. To specify an interposer type probe, set a value of 2. Zero indicates no probe installed. QpiHost.h provides equates that can be used such as PROBE_TYPE_NO_PROBE.
- **QpiConfigTxWci :** To configure what are the link width modes the link is allowed to operate.

6.2.27 Input Parameter Default Values

The following table is the default value for each QPIRC input option.

Table 1-1. Default Values

Option	Defaults
BusRatio/IoRatio/MmioIRatio	1
LegacyVgaSoc	0
MmioP2pDis	0
IsocAzaliaVc1En	0
DebugPrintLevel	0xF
AltRtid2S	0
QpiLinkSpeedMode	1
QpiLinkSpeed	1
ClusterOnDieEn	0
QpiLinkL0pEn	0



QpiLinkL1En	1
QpiL0rEn	1
QpiLbEn	0
IioUniphyDisable	0
QpiLinkCreditReduce	0
QpiConfigTxWci	0
QpiAdVNACreditArbThreshold	1
QpiCrcMode	0
QpiAdaptationEn	0
QpiAdaptationInParallel	0
QpiAdaptationSpeed	1
QpiCpuSktHotPlugEn	0
QpiCpuSktHotPlugTopology	0
QpiPortDisable	0
QpiLinkCreditReduce	0
QpiLinkSpeed	1
QpiProbeType	0
QpiConfigdTxWci	0

6.3 Output Parameters

Output parameters describe various aspects of the system QPI fabric. Some of the parameters are system wide and some are per socket/link.

6.3.1 CpuInfo

This is an array of MAX_CPU_SOCKETS which is configurable at build time. The array is indexed by socket id. It provides information about the system such as the sockets populated in the system, how they are connected to each other expressed in terms of LEP, the routes between sockets expressed in terms of topology tree and various resources allocated to each socket. The information provided will be sufficient enough to construct ACPI tables (such as SLIT), allocating the MBAR address for IOAPICs and allocating Bus/IO/MMIOL resources for each PCIe device in the system.

The following fields are part of this structure:

- **Valid:** Indicates if this socket is populated or not
- **SocId:** Socket id
- **LepInfo:** This is an array of MAX_QPI_PORTS which is 3 for these platforms. It contains the Link Exchange Parameters of the QPI links. The following fields are part of this structure:
 - **Valid:** Indicates if the link is operational, slow/full speed
 - **PeerSocId:** Socket Id of the peer socket
 - **PeerSocType:** Socket type of the peer socket (only SOCKET_TYPE_CPU is supported)
 - **PeerPort:** Port of the peer socket to which this link is connected
 - **DualLink:** Indicates if this link is part of pair of links connecting two sockets
- **TopologyInfo:** Describes the QPI topology of the system from this socket's perspective. Array element 0 always contains info about this socket (i.e root node) itself.
 - **Valid:** Indicates if this array element has valid information. An invalid array element indicates the termination of the topology tree
 - **SocId:** Socket id of this node
 - **SocType:** Socket type of this node
 - **ParentPort:** Port that connects this node to its parent node
 - **Hop:** Number of nodes between this node and root node including self
 - **ParentIndex:** Array index of this node's parent in the topology tree
- **CpuRes:** It contains the following fields that describe the system resource range assigned for each socket.
 - **BusBase, BusLimit:** Bus range decoded by this socket
 - **IoBase, IoLimit:** IO range decoded by this socket
 - **IoApicBase, IoApicLimit:** IOAPIC range decoded by this socket
 - **MmiolBase, MmiolLimit:** Low MMIO range decoded by this socket
 - **MmiohBase, MmiohLimit:** High MMIO range decoded by this socket

6.3.2 SysConfig

Indicates the QPI topology detected for the system. SYS_CONFIG structure in QpiSi.h defines all the possible values for this field.

6.3.3 WarningLog

Various warnings encountered during QPI initialization. Refer to “Warning Logs” section for more information.

6.3.4 OutLegacyVgaSoc

Indicates the socket ID which is configured to be the target of the legacy VGA region.

6.3.5 OutSocEn

This field indicates if value selected for **IsocEn** in sysHost Input structure by the platform BIOS is overridden or not. QPIRC will override an input option if the QPI topology doesn't allow the value selected for an option.

6.3.6 OutMesegEn

This field indicates if value selected for **MeSegEn** in sysHost Input structure by the platform BIOS is overridden or not. QPIRC will override an input option if the QPI topology doesn't allow the value selected for an option.

6.3.7 OutSocAzaliaVc1En

This field indicates if value selected for **IsocAzaliaVc1En** in QPIRC Input structure by the platform BIOS is overridden or not. QPIRC will override an input option if the QPI topology doesn't allow the value selected for an option.

6.3.8 OutClusterOnDieEn

This field indicates if the Cluster On Die feature is enabled.

6.3.9 OutMctpEn

This field indicates if the MCTP feature is enabled.

6.3.10 OutE2EParityEn

This field indicates if the end to end parity feature is enabled.

6.3.11 OutEarlySnoopEn

This field indicates if early snooping is enabled.

6.3.12 QpiCurrentLinkSpeedMode

This field indicates if the system is in Slow or Full speed operation.

6.3.13 OutQpiLinkSpeed

This field indicates link speed per link.

6.3.14 OutQpiLinkL0pEn

This field indicates if L0p is enabled or not.

6.3.15 OutQpiLinkL1En

This field indicates if L1 is enabled or not.

6.3.16 OutQpiLinkL0rEn

This field indicates if L0r is enabled or not.



6.3.17 OutlioUniphyDisable

This field indicates sockets for which the Uniphy is disabled.

6.3.18 OutQpiCrcMode

This field indicates CRC mode configured for the system.

6.3.19 RasInProgress

Set to TRUE when the QPIRC is called to handle a CPU/IOH event

6.3.20 RasEvtType

Indicates RAS event type. 0 is Online, 1 is Offline

6.3.21 RasSocId

Indicates RAS event on socket ID

6.3.22 ProgramNonBC

Flag to indicate if the Non BC RTA entries to be programmed

6.3.23 ProgramBC

Flag to indicate if the BC RTA entries to be programmed

6.4 Error/Status Codes

Output parameters describe various aspects of the system QPI fabric. Some of the parameters are system wide and some are per socket/link.

- [23:20] - Bitmask of CPU Sockets affected, 0xF - System Wide Error/Status
- [19:18] - Socket Type. 0: CPU Socket, 0x3 System Wide Error/Status
- [17:16] - Bit mask of QPI Links affected, 0 - Link N/A, 0x3 - System Wide Error/Status
- [15:08] - Error/Status Minor Code
- [07:00] - Error/Status Major Code

Each Major error code also has a Minor error code that describes the problem with more level of accuracy. Error/Status Major Code is also passed to OEM code by calling "OemCheckPoint" routine. This is typically display as POST code on IO port 0x80. All the errors are fatal and some of the errors might be transient. When these errors occur, the system will be reset.

Table 1-2. Major/Minor Error Codes

Major Error Code	Minor Error Code	Details
ERR_BOOT_MODE (0xD8)	MINOR_ERR_UNSUPPORTED_BOOT_MODE (0x01)	N/A
	MINOR_ERR_PBSP_CHKIN_FAILURE (0x02)	When system BSP tries to setup path for remote sockets or sends a Boot_Go command to remote socket in SetupSbspPathToAllSockets() or SyncUpPbspForReset(). If the remote socket(s) hasn't checked-in, assert; it is a fatal condition, this error will be logged. No retry. <i>RC Behavior: System Halt</i>
ERR_MINIMUM_PATH_SETUP (0xD9)	MINOR_ERR_ADD_SOCKET_TO_TOPOLOGY_TREE (0x01) MINOR_ERR_TOPOLOGY_TREE (0x02) MINOR_ERR_INTERNAL_DATA_STRUCTURE (0x03)	N/A N/A When SBSP tries to add this remote socket into system topology tree in SetupSbspPathToAllSockets(), there are some errors occur in the data structure. No retry.
	MINOR_ERR_NONCPU_PARENT_NODE (0x04)	<i>RC Behavior:</i>



	MINOR_ERR_INVALID_PORT_CONNECTION (0x05)	<p>The current Socket is not added to the tree.</p> <p>N/A</p> <p>When SBSP setups the boot path for the parent which is not directly connected to Legacy CPU in SetupSbspPathToAllSockets(). The Child is not an immediate neighbor of Parent. No retry.</p>
ERR_TOPOLOGY_DISCOVERY (0xDA)		N/A
ERR_SAD_SETUP (0xDB)	MINOR_ERR_INSUFFICIENT_RESOURCE_SPACE (0x01)	<p>N/A</p> <p>RC Behavior: System Halt</p>
ERR_UNSUPPORTED_TOPOLOGY (0xDC)	MINOR_ERR_INVALID_SOCKET_TYPE (0x01) MINOR_ERR_INVALID_CPU_SOCKET_ID (0x02) MINOR_ERR_CBO_COUNT_MISMATCH (0x03) MINOR_ERR_HA_COUNT_MISMATCH (0x04) MINOR_ERR_R3QPI_COUNT_MISMATCH (0x05) MINOR_ERR_SKU_MISMATCH (0x06) MINOR_ERR_LEGACY_PCH_MISMATCH (0x07) MINOR_ERR_SBO_COUNT_MISMATCH (0x08)	<p>N/A</p> <p>RC Behavior: System Halt</p>
ERR_FULL_SPEED_TRANSITION (0xDD)	MINOR_ERR_LINK_SPEED_UNSUPPORTED (0x01) MINOR_ERR_MAX_CLK_RATIO_UNSUPPORTED (0x02) MINOR_ERR_QPI_ELEC_PARAM_NOT_FOUND (0x03)	<p>N/A</p> <p>SBSP cannot find QPI TXEQ Parameters for this link in GetSocketLinkEparams(). No retry.</p> <p>RC Behavior: System Halt</p>

Table 1-3. Major/Minor Status Codes

Major Status Code	Minor Status Code
STS_DATA_STRUCTURE_INIT (0xA0)	N/A
STS_COLLECT_EARLY_SYSTEM_INFO (0xA1)	N/A
STS_OPEN_S BSP_CONFIG_ACCESS (0xA2)	N/A
STS_SETUP_MINIMUM_PATH (0xA3)	MINOR_STS_ADD_SOCKET_TO_MIN_PATH_TREE (0x01) MINOR_STS_COLLECT_LEP (0x02) MINOR_STS_CHK_PBSP_CHKIN (0x03) MINOR_STS_SET_PBSP_BOOT_PATH (0x04) MINOR_STS_SET_S BSP_CONFIG_PATH (0x05)
STS_OPEN_PBSP_CONFIG_ACCESS (0xA4)	N/A
STS_EARLY_RESET_CONFIG (0xA5)	N/A
STS_PBSP_SYNC_UP (0xA6)	MINOR_STS_ISSUE_BOOT_GO (0x01)
STS_TOPOLOGY_DISCOVERY (0xA7)	MINOR_STS_CHK_TOPOLOGY (0x01) MINOR_STS_DETECT_RING (0x02) MINOR_STS_CONSTRUCT_TOPOLOGY_TREE (0x03) MINOR_STS_CALCULATE_ROUTE (0x04)
STS_PROGRAM_FINAL_ROUTE (0xA8)	N/A
STS_PROGRAM_FINAL_IO_SAD (0xA9)	MINOR_STS_ALLOCATE_CPU_RESOURCE (0x01) MINOR_STS_FILL_SAD_TGTLST (0x02) MINOR_STS_PROGRAM_CPU_SAD_ENTRIES (0x03)
STS_PROTOCOL_LAYER_SETTING (0xAA)	MINOR_STS_PROGRAM_RTID (0x01) MINOR_STS_PROGRAM_RING_CRDT (0x02) MINOR_STS_PROGRAM_HA_CRDT (0x03) MINOR_STS_PROGRAM_HTRBT (0x04) MINOR_STS_QPI_MISC_SETUP (0x05)
STS_FULL_SPEED_TRANSITION (0xAB)	N/A
STS_PHY_LAYER_SETTING (0xAC)	N/A
STS_LINK_LAYER_SETTING (0xAD)	N/A
STS_SYSTEM_COHERENCY_SETUP (0xAE)	N/A
STS_QPI_COMPLETE (0xAF)	N/A

6.5 Warning Logs

QPIRC logs four-byte wide warning data in the output structure. The format of the warning is as follows:

- [15:12] - Bitmask of CPU Sockets affected, 0xF - System wide Warning
- [11:10] - Socket Type. 0 - CPU Socket, 0x3 - System Wide Warning



[09:08] - Bit mask of QPI Links affected. 0x3 - System Wide Warning, 0 – Link N/A

[07:00] - Warning Code

The warning log size in the output structure is pre-determined. So whenever the log overflows, the last entry of the warning log will indicate if there is any overflow.

Table 1-4. Warning Codes

Warning Codes	Details
WARN_LINK_SLOW_SPEED_MODE (0x01)	After warm-reset, QPIRC will check the link speed and this warning message will be logged if the system is operating in slow speed mode in QpiPhyLinkAfterWarmReset()
WARN_UNSUPPORTED_LINK_SPEED (0x02)	If requested speed is unsupported, this warning message will be logged. In SelectSupportedLinkSpeed().
WARN_LINK_FAILURE (0x03)	N/A
WARN_PER_LINK_OPTION_MISMATCH (0x04)	Normalize per link options. If the peers don't match their options, they will be normalized and this warning message will be logged In NormalizeLinkOptions()
WARN_MMIOH_BASE_UNSUPPORTED (0x09)	N/A
WARN_MMIOH_SIZE_UNSUPPORTED (0x0A)	If user sets MMIOH size requested is 0, QPIRC force to 2G and this warning will be logged in AllocatelioResources()
WARN_RESOURCE_NOT_REQUESTED_FOR_CPU (0x0C)	If user selects 0 for IO and MMIOH resource ratio on legacy socket, this warning will be logged in CalculateResourceRatio()
WARN_RESOURCE_REQUEST_NOT_MET (0x0D)	Log a warning if the platform requested resource is more than what is available in CalculateResourceRatio().
WARN_VGA_TARGET_SOC_NOT_PRESENT (0x0E)	If legacy VGA socket is out of the max socket range or not a valid socket, this warning message will be logged in PrimeHostStructure().

6.6 Porting Guidelines

The following files need to be ported to use the QPIRC in an OEM BIOS code base:

6.6.1 QpiHooks.c

These following routines require porting:

- o OemDebugPrintQpi(): Simple printf() style implementation that can print different levels of debug messages.
- o OemWaitTimeForPSBP(): Can be used to specify a max wait time for the PSBP to check in.
- o OemQpiGetEparams(): This function must be ported to return the full-speed electrical parameters of the link to the calling routine. A simple example is provided in the reference code, but can be expanded to detect and support multiple platforms. This enables a single QPIRC binary to support multiple platforms.
- o OemQpiGetHalfSpeedEParams() This function must be ported to return the half-speed electrical parameters of the link to the calling routine. A simple example is provided in the reference code, but can be expanded to detect and support multiple platforms. This enables a single QPIRC binary to support multiple platforms.
- o OemCheckCpuPartsChangeSwap(): This function should detect if the CPU parts have changed or swapped for the platform since last boot. Used to avoid rerunning adaptation if the parts parts have not changed/swapped.
- o OemGetAdaptedEqSettings(): This function should return the adapted EQ values found for the platform.

6.6.2 QpiPlatformEparam.c

This file contains all the electrical parameters for all the sockets, links, and speeds. These values typically come from the board designers.

7 Memory Reference Code

7.1 Basic Terminologies and Assumptions

Throughout the MRC the following terms are used:

- CPU0, P0, Socket 0, etc.
This is the processor package that resides on a programmable PCI bus Number. The MRC expects as an input which CPUs are installed on a given system.
- Node0.
The MRC passes the memory node # to the routines to address/access the PCI config space of the Node it's trying to access and the routine reads the programmed bus number for that socket.
- DIMM0, D0, etc
This is the DIMM slot at the far end of a memory channel away from the processor socket. DIMM modules in a channel must be populated starting from this slot first.
- ..._P0C0D0, P1C2D1, etc.
These are used in MEMPLATFORM.H file to represent the locations of DIMM slots. For example, ..._P1C0D1 means CPU1 or Processor 1 (see above), Channel 0, DIMM slot 1.

7.2 MRC Entry Point

The 'C' prototype for the MRC entry point is

```
UINT32 MemStart (PSYSHOST host);
```

The parameter is a pointer to a structure on the stack that contains input setup data, system variable data and nvram data. The MRC will also use the sysVar and sysNvram data as output.

```
typedef struct sysHost {
    const struct sysSetup  setup; // constant input data
    struct sysVar          var;   // variable, volatile data
    struct sysNvram        nvram; // variable, non-volatile data for S3, etc...
    struct prevBootErrInfo prevBootErr; // Collect UNCORE MC Bank information.
#ifdef BDAT_SUPPORT
    BDAT_STRUCTURE bdat;
#endif // BDAT_SUPPORT
    UINT32 cpuFlows; // Which CPU flows to take in simulation.
    UINT32 qpiFlows; // Which QPI flows to take in simulation.
    UINT32 memFlows; // Which Memory flows to take in simulation.
}
SYSHOST, *PSYSHOST;
```

Please refer to the sample code of routine `GetMemSetupOptions()` in the file `MemSetup.c` as an example of how these input parameters are prepared before calling the entry point `MrcStart()`.

7.3 Setup Options and Input Parameters

The memory controller provides a highly flexible memory subsystem, allowing a wide variety of configuration options with respect to interleave schemes, channel sparing, mirroring, etc. The configuration options that require user input will come from platform-specific implementation, such as the regular BIOS Setup. The BIOS evaluates these setup options by invoking 32-bit procedures to read NVRAM contents. This technique is not practical within the MRC due to the 32-bit execution environment and separate link phase of the BIOS build.

The solution is to evaluate all required setup options before calling into the MRC. The main BIOS module has programmatic access to NVRAM access procedures and capable of calling them as needed. After initializing CAR mode and establishing a stack, the main BIOS can invoke NVRAM procedures, for example, to read setup options and initialize a predefined data structure on the stack. The MRC utilizes the common SYSHOST data structures for all input and output parameters.

Typical platform specific settings and hooks for the BIOS reference code are contained in the platform directory's subdirectories, with MRC specific files pre-pended with the string "Mem". An OEM porting the BIOS reference code for a



specific platform should carefully examine all the files under the platform directory to ensure proper operation of the MRC on a given platform.

Data structures common to the MRC and other parts of the BIOS reference code are defined in SysHost.h file, and the MRC specific data structures are defined in file MemHost.H. These header files are consumed by the 32-bit MRC source. The main BIOS is also encouraged to use MRC equates defined in this include file.

7.3.1 Setup Structure Initialization

The main BIOS is first required to create a stack frame for the **sysSetup** structure and initialize the contents using either constant or runtime data. **Sample code is located in the files CommonSetup.c and MemSetup.c.** The main BIOS developer can modify this procedure as needed to read setup options using native NVRAM access routines.

The format of the sysSetup structure may vary with each RC release. The individual BIOS developer must review this structure to ensure that all fields are properly initialized. The individual field definitions for Intel® Xeon® xxxx Platform MRC xxxx are described below:

7.3.2 sysSetup

The sysSetup structure contains the “selectable” input parameters used by the BIOS Reference Code. Memory specific parameters are contained in the memSetup structure that is included in the sysSetup structure. The values to the fields in this structures can be customized to create the MRC behavior desired by the platform BIOS.

7.3.2.1 options

This is a DWORD bit map field with the following options:

7.3.2.1.1 TEMPHIGH_EN (BIT0)

When set, enables support for 95 degree DIMMs. These DIMMs will operate at 85 degrees if this bit is not set. This is needed because some platforms may not be able to handle the extra cooling required for 95 degree DIMMs. If this bit is not set, MRC will not set the SRT bit of MR2 register of DIMMs.

7.3.2.1.2 ATTEMPT_FAST_BOOT_COLD (BIT1)

When set, enables the reference code to attempt the fast cold boot path.

7.3.2.1.3 PDWN_SR_CKE_MODE (BIT2)

Set to enable CKE output tri-state control during self-refresh. When clear, the CKE is not tri-stated during SR. UDIMM configurations must have this bit cleared. When set, the CKE is tri-stated during register clock off power down self-refresh. It is recommended that BIOS should set this bit to 1 for RDIMM configuration.

7.3.2.1.4 OPP_SELF_REF_EN (BIT3)

Set to enable for Opportunistic Self-Refresh

7.3.2.1.5 MDLL_SHUT_DOWN_EN (BIT4)

Set to enable for Master DLL shut-down. When set, all MDLL's (command/control and data) will be Shut-down.

7.3.2.1.6 PAGE_POLICY (BIT5)

Clear for open page, set for closed page. Setting this bit will cause MRC to set CLOSED_PAGE and AUTOPRECHARGE bits of MC_CONTROL register in EDS. It also affects the DIMM static power calculation for Thermal Throttling support.

7.3.2.1.7 ALLOW2XREF_EN (BIT6)

Set to enable the MRC to use 2x refresh if needed for DIMMs that support the extended operating temperature range (95 degree DIMMs). If set and TEMPHIGH_EN is also set, tREFI_8 timing of the channel will be reduced to half for 2x refresh rate. Note that all DIMMs on the same channel must support extended temperature range and x2 refresh before MRC honors the x2 refresh.

7.3.2.1.8 MULTI_THREAD_MRC_EN (BIT7)



When set, memory configuration will parallelized where appropriate, utilizing socket BSP for each present socket to assist in memory initialization. When cleared, memory initialization will be completely serialized from legacy socket BSP.

7.3.2.1.9 ADAPTIVE_PAGE_EN (BIT8)

Set to enable Adaptive Page Close. Refer to ADAPTIVEPAGECLOSE register field in EDS.

7.3.2.1.10 CMD_CLK_TRAINING_EN (BIT9)

When set, command clock training will be completed during memory initialization. When clear, command clock will not be trained.

7.3.2.1.11 SCRAMBLE_EN (BIT10)

Set to enable memory data scrambling via MCSCRAMBLE register. This bit should always be set to 1 except for LA-based debugging.

7.3.2.1.12 BANK_XOR_EN (BIT11)

Enable bank xor address mapping.

7.3.2.1.13 RESERVED (BIT12) DDR_RESET_LOOP (BIT13)

Set to enable reset loop testing after memory training. Can be used while gathering memory margining information during product qualification.

7.3.2.1.14 NUMA_AWARE (BIT14)

When set, memory initialization will optimize for Non-Uniform Memory Access by enabling socket interleaving where possible. When clear, socket interleaving will be disabled.

7.3.2.1.15 DISABLE_WMM_OPP_READ (BIT15)

Set to disable issuing read commands opportunistically during WMM .

7.3.2.1.16 Reserved (BIT16)

7.3.2.1.17 ECC_CHECK_EN (BIT17)

Set to allow MRC to enable ECC Checking.

7.3.2.1.18 ECC_MIX_EN (BIT18)

Set to allow MRC to boot with ECC enabled when a mixed configuration of ECC and non-ECC memory are installed in the system. MRC will disable all channels with non-ECC memory in this configuration and boot with ECC enabled, assuming there is at least one channel with ECC only memory present.

7.3.2.1.19 BALANCED_4WAY_EN (BIT19)

Set to enable a more optimal way of interleaving non-NUMA DP platforms that have a 2-1-1 channel configuration (2 DIMM on one channel and 1 DIMM on each of the other two channels). If the largest 2 channels are on different sockets and the sum of the 2 largest channels is more than the next 4 channels and the planned interleave is 6, force a 4-way interleave to make the performance more symmetric

7.3.2.1.20 CA_PARITY_EN (BIT20)

Set to enable Command/Address parity.

7.3.2.1.21 SPLIT_BELOW_4GB_EN (BIT21)

Set to enable the allocation of memory below 4GB to both processor sockets when in NUMA mode. This can be enabled for performance reasons under certain configurations. Some operating systems require this to be disabled. The default should be disabled.

7.3.2.1.22 RESERVED (BIT22)

This bit is reserved for future use.



7.3.2.1.23 RAS_TO_INDP_EN (BIT23)

When set, switches from lockstep or mirror mode to independent channel mode and log a warning when memory is found present on channel 2 for a SKT R3 processor. If cleared, the MRC will ignore DIMMs present on channel 2.

7.3.2.1.24 RESERVED (BIT24)

Not used.

7.3.2.1.25 MARGIN_RANKS_EN (BIT25)

Set to enable DIMM margin check for HVM testing purposes. The read/write timing and voltage margins are stored in the mrcHost structure and output via serial debug messages.

7.3.2.1.26 MEM_OVERRIDE_EN (BIT26) (*if XMP_SUPPORT is defined*)

Set to enable memory initialization using the input parameters to override some memory configuration inputs.

7.3.2.1.27 DRAMDLL_OFF_PD_EN (BIT27)

Set to enable DRAM OFF Powerdown Slow Mode in DIMM when performing self refresh.

7.3.2.1.28 MEMORY_TEST_EN (BIT28)

Set to run memory initialization built in self test(BIST) on normal boot paths. When clear, Memory BIST will not be executed.

7.3.2.1.29 MEMORY_TEST_FAST_BOOT_EN (BIT29)

When set, enables the memory test when going through a fast boot path. When clear, memory test will be skipped on a fast boot path.

7.3.2.1.30 ATTEMPT_FAST_BOOT (BIT30)

When set, memory initialization will attempt to take a fast boot path if the NVRAM structure is good and the memory config hasn't changed. When clear, memory initialization will use normal boot path. For example, on a warm boot, this will take the "fast warm" path through MRC which attempts to make it as close as possible to the S3 path.

7.3.2.1.31 SW_MEMORY_TEST_EN (BIT31)

Set to enable memory initialization performing a software memory test. When clear, software memory test will not be performed.

7.3.2.1.32 FASTBOOT_MASK (BIT30 | BIT1)

Used to detect if above option fields have changed since previously saved cold boot data.

7.3.2.2 optionsExt

7.3.2.2.1 EARLY_CMD_CLK_TRAINING_EN (BIT0)

When set, enable early command clock training.

7.3.2.2.2 RD_VREF_EN (BIT1)

When set, enable Read Vref training.

7.3.2.2.3 WR_VREF_EN (BIT2)

When set, enable Write Vref training.

7.3.2.2.4 XOVER_EN (BIT3)

When set, enable Crossover training.

7.3.2.2.5 SENSE_EN (BIT4)

When set, enable Sense Amp training.

7.3.2.2.6 PDA_EN (BIT5)

When set, enable PDA mode (MR3 bit4).

7.3.2.2.7 TURNAROUND_OPT_EN (BIT6)

When set, memory initialization code will optimize the turnaround timing based on training results (TCRWP, TCOTHP, TCOTHP2, TCLRDP registers).

7.3.2.2.8 ROUND_TRIP_LATENCY_EN (BIT7)

When set, memory initialization code will optimize the Round Trip timing and I/O latency (ROUNDTRIPx, IOLATENCYx registers).

7.3.2.2.9 ALLOW_CORRECTABLE_ERROR (BIT8)

When set, allow correctable error during MRC training (do not enforce memory mapout/disabling on SBE error).

7.3.2.2.10 EARLY_CTL_CLK_TRAINING_EN (BIT9)

When set, enable Early CTL-CLK training.

7.3.2.2.11 TX_EQ_EN (BIT10)

When set, enable TX_EQ.

7.3.2.2.12 IMODE_EN (BIT11)

When set, enable i mode training.

7.3.2.2.13 CMD_NORMAL_EN (BIT12)

When set, enable CMD normalization.

7.3.2.2.14 PER_BIT_MARGINS (BIT13)

When set, enabled Per Bit Margin data.

7.3.2.2.15 DUTY_CYCLE_EN (BIT14)

When set, enable Duty Cycle training.

7.3.2.2.16 RCOMP_TYPE (BIT15)

Set to enable RCOMP (bit 22) of PCU MISC CONFIG (MAILBOX_BIOS_CMD_WRITE_PCU_MISC_CONFIG)

7.3.2.2.17 RESERVED (BIT16 – BIT31)

Reserved for future use.

7.3.2.3 bclk

This is a byte sized field that selects the bclk frequency (e.g. 100, 133, 145, etc). Selecting a value of 0 results in the MRC automatically selecting the bclk value..

7.3.2.4 enforcePOR

This is a byte sized field that contains information about memory DIMM population rule. 0 - Enforce POR, 1 – Enforce, Stretch goal, 2 – No Limit. Default – Enforce POR.

7.3.2.5 RASmode

This is a byte sized field that contains information about which RAS mode is desired. The options are a bit field where bits 0, 1, and 2, represent CH_MIRROR, CH_LOCKSTEP and RK_SPARE respectively. If no bits are set, then CH_INDEPENDENT is selected. The MRC will verify if the desired mode is possible for the given memory population. CH_INDEPENDENT is the default for invalid RAS selections.



7.3.2.6 RASmodeEx

This is a byte sized field that contains information about which additional RAS features are desired. The MRC will verify if the desired mode is possible for the given system and force the option to disable and log a warning if appropriate.

```
//
// Host extended RAS Modes
//
#define DEVTAGGING_EN      BIT0
#define DMNDSCRIB_EN      BIT1
#define PTRLSRIB_EN      BIT2
#define A7_MODE_EN      BIT4
```

7.3.2.7 ddrFreqLimit

Forces a maximum DDR operation frequency regardless of the common tCK detected via SPD of all populated DIMMs. Options are AUTO, DDR_800, DDR_1000, DDR_1067, DDR_1200, DDR_1333, DDR_1400, DDR_1600, DDR_1800, DDR_1867, DDR_2000, DDR_2133, DDR_2200, DDR_2400, DDR_2600, DDR_2800, DDR_2933, DDR_3000, and DDR_3200. AUTO will set the fastest common speed as determined from the SPD information of all populated DIMMs. A ddrFreqLimit value higher than the fastest common speed detected will be ignored.

7.3.2.8 SocketInter

Nodes (or processor) interleave setting. Valid options are 1, 2, or 4 way interleave. Default 2 ways interleave.

7.3.2.9 chInter

Channels interleave setting. Valid options are 1, 2, 3, or 4 way interleave. Other values defaults to 4 ways interleave.

7.3.2.10 rankInter

Rank interleaves setting. Valid options are 1, 2, 4, or 8 way interleave. Other values defaults to 8 ways interleave.

7.3.2.11 dimmTypeSupport

Type of DDR DIMM supported by this platform. Valid options are 0 – RDIMM, 1 - UDIMM, 2 – Both. The MRC will enforce population rules based on this parameter. Even if “Both” is selected the MRC will not allow RDIMMs and UDIMMs to be installed at the same time, but will automatically choose the first type detected.

7.3.2.12 vrefStepSize

Amount to inc/dec Vref DCP register value for each step of the Vref training algorithm.

7.3.2.13 vrefAbsMaxSteps

Max number of steps for DIMM Vref margin +/-10% of Vddq

7.3.2.14 vrefOpLimitSteps

Max number of steps for DIMM Vref margin +/-5% of Vddq/2

7.3.2.15 pdwnCkMode

Selects the Power Down Clock Mode for UDIMM or RDIMM. This field defines how CK and CK# are turned off during Self Refresh. See EDS for possible values. ckeThrottling Selects the CKE Power management mode. See EDS for possible values.

7.3.2.16 MemPwrSave

Selects the memory power saving mode. Depends on the selected mode, the Power Down clock mode, CKE, and IBT are initialized accordingly.



7.3.2.17 ckeThrottling

Configures CKE Throttling. Valid options are 0 – “OFF”, 1 - “APD On”, 2 - “PPDF On”, 3 - “PPDS On”, 4 - “APD On, PPDF On”, 5 - “APD On, PPDS On”.

7.3.2.18 olttPeakBWLIMITPercent

$(\text{value}/100) * 255 / \text{max number of dimms per channel} = \text{DIMM_TEMP_THRT_LMT THRT_HI}$.

7.3.2.19 thermalThrottlingOptions

Bit fields for selecting the enabling Thermal Throttling Modes. Bits exist for Open Loop Thermal Throttling (OLTT) and Closed Loop Thermal Throttling (CLTT) as well as mem_hot output generation logic and Enabling external MEM_HOT sensing logic.

7.3.2.20 dramraplen

Selects whether the reference code will initialize and enable DRAM RAPL.

7.3.2.21 dramraplbwlmittf

Allows custom tuning of BW_LIMIT_TF when DRAM RAPL is enabled.

7.3.2.22 dramraplRefreshBase

Allows custom tuning of Power scaling by Refresh rate in units of 0.1x when DRAM RAPL is enabled

7.3.2.23 dramMaint

When set, enables DRAM maintenance settings.

7.3.2.24 dramraplExtendedRange

When set, enable extended range for DRAM RAPL.

~~7.3.2.24~~ 7.3.2.25 dramMaintTRRMode

Select mode A or B for TRR mode.

~~7.3.2.25~~ 7.3.2.26 dramMaintMode

Select pTRR or TRR mode.

~~7.3.2.26~~ 7.3.2.27 electricalThrottling

Controls enabling of electrical throttling. When value is ET_DISABLE, electrical throttling is disabled. When value is ET_ENABLE, electrical throttling is forced enabled. When value is ET_AUTO, memory initialization code will select enable/disable based on memory topography.

~~7.3.2.27~~ 7.3.2.28 altitude

Selects the system altitude for memory thermal throttling calculations.

~~7.3.2.28~~ 7.3.2.29 forceRankMult

Selects LRDIMM rank multiplication factor.

~~7.3.2.29~~ 7.3.2.30 lrdimmModuleDelay

LRDIMM Module Delay. 0 – Delay auto, 1 – disable. Default 0

~~7.3.2.30~~ 7.3.2.31 lrdimm_x4asx8

Option to run x4 LRDIMM as a x8 LRDIMM for power savings.



~~7.3.2.31~~ 7.3.2.32 **rxVrefTraining**

Enables the Rx Vref advanced training step.

~~7.3.2.32~~ 7.3.2.33 **perBitDeSkew**

Enables the Per Bit Deskew advanced training step.

~~7.3.2.33~~ 7.3.2.34 **iotMemBufferRsvtn**

Reserves a block of memory as IOT trace buffer for each CPU socket.

~~7.3.2.34~~ 7.3.2.35 **setMCODT**

Select MCODET ("0" = 50 ohms, "1" = 100 ohms)

~~7.3.2.35~~ 7.3.2.36 **rmtPatternLength (if MARGIN_CHECK is defined)**

Selects the pattern length for the rank margin tool. One pattern equals the size of the write data buffer.

~~7.3.2.36~~ 7.3.2.37 **patrolScrubDuration**

Specifies the number of hours requested for patrol scrub to scrub all system memory.

~~7.3.2.37~~ 7.3.2.38 **spareErrTh**

Spare error threshold. Contains number of correctable ECC errors required before triggering a SMI event. This value will be programmed into the cor_err_th_0 and cor_err_th_1 fields of the CORRERRTHRSHLD registers for all channels and all sockets. Default value is 0x7FFF. The max value is 0x7FFF.

~~7.3.2.38~~ 7.3.2.39 **leakyBktLo**

Leaky bucket low mask position.

~~7.3.2.39~~ 7.3.2.40 **leakyBktHi**

Leaky bucket High mask position.

~~7.3.2.40~~ 7.3.2.41 **SpareRanks**

Number of ranks which could be spared.

~~7.3.2.41~~ 7.3.2.42 **MemTestLoops**

Number of MemTest loops to run. Default to 0 (AUTO) which means only 1 test will run if MemTest is enabled. This option is not valid if MemTest is disabled (BYPASS_HWMEMTEST_EN=1).

~~7.3.2.42~~ 7.3.2.43 **scrambleSeedLow**

Setting low 16 bits of the data scrambling seed.

~~7.3.2.43~~ 7.3.2.44 **scrambleSeedHigh**

Setting high 16 bits of the data scrambling seed.

~~7.3.2.44~~ 7.3.2.45 **ADREn**

Enables support for ADR.

~~7.3.2.45~~ 7.3.2.46 **eraseArmNVDIMMs (if MEM_NVDIMM_EN is defined)**

When set, MRC will perform NVDIMM erase and Arm operations on this boot.

~~7.3.2.46~~ 7.3.2.47 **check_pm_sts**

Enables use of the PCH_PM_STS register as a recovery indicator.

~~7.3.2.47~~ 7.3.2.48 **check_platform_detect**

Enables use of the PlatformDetectADR function as a recovery indicator.

~~7.3.2.48~~ 7.3.2.49 **mcBgfThreshold**

The HA to MC BGF threshold is used for scheduling MC request in bypass condition. Valid values are from 0 to 15. Default value is 4.

~~7.3.2.49~~ 7.3.2.50 **lockstepEnableX4**

When clear, overrides other lockstep configuration options when x4 DIMMS are detected.

~~7.3.2.50~~ 7.3.2.51 **NormOpplIntvl**

Normal Operation Interval

~~7.3.2.51~~ 7.3.2.52 **NumSparTrans**

Number of sparing transactions interval.

~~7.3.2.52~~ 7.3.2.53 **SpdSmbSpeed**

Select DDR4 SMB Speed.

~~7.3.2.53~~ 7.3.2.54 **oemHooks**

See oemHooksSetup.

~~7.3.2.54~~ 7.3.2.55 **Per Memory Node Options**

See ddrIMCSetup.

~~7.3.2.55~~ 7.3.2.56 **inputMemTime (XMP_Support only)**

See memTiming structure.

~~7.3.2.56~~ 7.3.2.57 **XMPChecksum (XMP_Support only)**

UINT16 XMPChecksum per [MAX_NODE][MAX_CH][MAX_DIMM].

~~7.3.2.57~~ 7.3.2.58 **meRequestedSize (if ME_SUPPORT_FLAG is defined)**

Requested size for ME UMA in 1MB units.

~~7.3.2.58~~ 7.3.2.59 **phaseShedding**

Enable/Disable memory VR phase shedding feature. 0 – Disable, 1 – Enable, Default - Enable

~~7.3.2.59~~ 7.3.2.60 **DimmTempStatValue**

Dimm Temperature value.

~~7.3.2.60~~ 7.3.2.61 **customRefreshRate**

Allows custom tuning of Refresh rate from 2.0x to 4.0x in units of 0.1x.

~~7.3.2.61~~ 7.3.2.62 **DramMaintenanceTest**

When set, enable dram maintenance test



~~7.3.2.62~~ 7.3.2.63 **DMTDir**

For use with dram maintenance test: changing test direction UP or DOWN.

~~7.3.2.63~~ 7.3.2.64 **DMTInv**

For use with dram maintenance test: pattern inversion.

~~7.3.2.64~~ 7.3.2.65 **DMTRep**

For use with dram maintenance test: fast hammer number of repetitions.

~~7.3.2.65~~ 7.3.2.66 **DMTIter**

For use with dram maintenance test: number of hammer iterations per row.

~~7.3.2.66~~ 7.3.2.67 **DMTSwizzleEn**

For use with dram maintenance test: rowSwizzleType (ROW_SWIZZLE_3XOR1_3XOR2 or ROW_SWIZZLE_NONE)

~~7.3.2.67~~ 7.3.2.68 **DMRefreshEn**

For use with dram maintenance test: when set, enable refresh.

~~7.3.2.68~~ 7.3.2.69 **dllResetTestLoops**

Number of times to loop through RMT to test the DLL Reset

7.3.3 **ddrChannelSetup**

This structure contains parameters specific to each DDR3 channel within a socket. There is a copy of this structure for each channel on each socket.

7.3.3.1 **enabled**

Used to enable or disable a channel. Non-zero: Enable the channel, 0: Disable the channel

7.3.3.2 **options**

Reserved

7.3.3.3 **numDimmSlots**

Number of DIMM slots per channel actually present on the board. Valid options are 1, 2 or 3. MAX_DIMM is defined in mrcplatform.h as a build-time parameter. This option can be no larger than MAX_DIMM. It overrides MAX_DIMM when it is smaller.

7.3.3.4 **batterybacked**

Selects whether channel has a battery backup.

7.3.3.5 **rankmask**

Rank map-out based on bit set indication.

7.3.3.6 **vrefDcp**

Defines override of DCP SMBus device properties

7.3.3.7 **dimmlist**

This structure contains parameters specific to each DIMM within a channel

7.3.3.7.1 ddrDimmSetup

This sub-structure contains parameters specific to DIMMs in a channel. There is a copy of this sub-structure for each DIMM slot in a channel

7.3.3.7.2 mapOut (Byte)

If set to a non-zero value by the caller, the DIMM at this slot will be mapped out (disabled). Note that all DIMMs behind this DIMM on the same channel will also be mapped out per DIMM population rule, so the caller should set the corresponding mapOut fields of all DIMMs in a channel correctly to reflect this rule. MRC will enforce this rule if the input is incorrect.

7.4 DIMM Population Rules

The following table contains DIMM population rules enforced by the MRC. This table should be used in conjunction with the tables of supported DDR configurations to determine DIMMs supported and population rules. The Supported DDR Configuration POR table can be found in the Haswell EDS.

7.5 Platform Porting and Internal Hooks

The MRC integrates support for platform-specific controls during build and execution time. Platform settings such as the number of DIMMs per channel and chipset GPIO assignments are essential for proper memory detection and initialization. The MRC configuration that is expected to change per platform has been isolated in two MRC files: MemPlatform.H and MemHooks.C. A sample file MemSetup.C is provided to illustrate setting up the MRC input parameters and calling MRC's entry point. Other files in the Refcode\Platform subdirectories must also be ported for a given platform. **Other files in the RefCode directory and subdirectories should not be modified.** By default, these files implement settings compatible with the Intel customer reference board and the Intel Platform Design Guide.

7.5.1 MemPlatform.H file

This file defines the platform and BIOS-specific settings that are substituted into core MRC files during build time.

7.5.2 MemHooks.c file

This file defines platform and BIOS-specific hooks that the MRC will call during execution time. Currently, the MRC provides hooks for checkpoint display and memory vref control. The MemHooks.C file describes the execution context and interface of each hook.

Additional hooks will be added during MRC development as the need arises. If any MRC customer requires additional hooks for a specific application, they are encouraged to forward a request to their Intel representative. The request should include a description of the application with specific changes to MRC source files. The request should be submitted as early as possible to provide adequate time for validation.

7.6 POST and Error Codes

7.6.1 POST Codes

A POST code is a hexadecimal value indicating a specific point in the flow of code execution and is a useful tool for debugging. A POST code is sent to IO port 80h (or a different IO port defined by input parameter "debugPort"), stored into a 32-bit scratchpad register at BIOSNONSTICKYSCRATCHPAD7 (Uncore Bus, Device 11, Function 3, offset 0x7C) in following format:

[31:24] Major POST Code
[23:16] Minor POST Code
[15:0] Intel implementation specific

The POST code handler OutputCheckpoint() is part of the Checkpoint.c file which can be customized by OEMs to meet their specific needs. For serial MRC execution it is stored in SBSP uncore. For parallel MRC execution it is stored in all populated sockets.

The table below is a list of POST codes used by MRC.

Table 7-1 MRC POST Code

(MemPlatform.h)



POST Code Nomenclature (As defined in MemPlatform header file)	MajorCode	MinorCode	SubCode
STS_DIMM_DETECT: Detect DIMM population	0B0h		
#define SUB_DIMM_DETECT		0x00	
#define SUB_DISPATCH_SLAVES		0x01	
#define SUB_INIT_SMB		0x02	
#define SUB_GET_SLAVE_DATA		0x03	
STS_CLOCK_INIT: Set DDR frequency	0B1h		
#define SUB_SEND_FREQ		0x01	
#define SUB_SEND_STATUS		0x02	
#define SUB_SET_FREQ		0x03	
#define SUB_SET_CLK_VDD		0x04	
#define SUB_XMP_INIT		0x05	
STS_SPD_DATA: Gather remaining SPD data	0B2h		
#define SUB_SEND_DATA		0x01	
STS_GLOBAL_EARLY: Program registers on the memory controller level	0B3h		
#define SUB_PROG_TIMINGS		0x01	
STS_RANK_DETECT: Evaluate RAS modes and save rank information	0B4h		
#define SUB_SEND_DATA_SPD		0x01	
#define DATA_TYPE_VAR			0x01
#define DATA_TYPE_NVRAM_COMMON			0x02
#define DATA_TYPE_NVRAM_DATA			0x03
STS_CHANNEL_EARLY: Program registers on the channel level	0B5h		
STS_DDRIO_INIT: DDRIO Initialization sequence	0B6h		
STS_DDR_CHANNEL_TRAINING: Train DDR	0B7h		
#define SUB_WR_FLY_BY_TRAINING		0x01	
#define SUB_WR_FLY_BY_EVALUATE		0x02	
#define SUB_ENABLE_PARITY		0x03	
#define STS_RMT		0x10	
#define SUB_LATE_CONFIG			0x01
#define SUB_INIT_THROTTLING			0x02
#define STS_MINOR_END_TRAINING		0x11	
#define STS_MINOR_NORMAL_MODE		0x12	
#define SUB_DATA_SYNC			0x01
STS_INIT_THROTTLING: Initialize CLTT/OLTT	0B8h		
STS_MEMBIST: Hardware memory test and init	0B9h		
STS_MEMINIT: Execute memory init	0BAh		
STS_DDR_MEMMAP: Program memory map and interleaving	0BBh		
#define SUB_SAD_INTERLEAVE		0x01	
#define SUB_TAD_INTERLEAVE		0x02	
#define SUB_SAD_NONINTER		0x03	
#define SUB_TAD_NONINTER		0x04	
#define SUB_RANK_INTER		0x05	
#define SUB_WRITE_SAD		0x06	
#define SUB_WRITE_TAD		0x07	
#define SUB_WRITE_RIR		0x08	
#define SUB_WRITE_SPARE_RANK		10	
STS_RAS_CONFIG: Program RAS configuration	0BCh		
#define SUB_MIRROR_MODE		0x01	
#define SUB_SPARE_MODE		0x02	
#define SUB_DEVICE_TAG		0x03	
#define SUB_ERR_THRESH		0x04	
STS_GET_MARGINS:	0xBD		
STS_MRC_DONE: MRC is done	0BFh		

7.6.2 Post Code Stall

The OutputCheckpoint() routine checks BIOS SCRATCHPAD6 for a matching major/minor checkpoint code. If the checkpoint codes match the BIOS stalls waiting for a change in code. This can be used by EV to pause execution at specific training steps. Validation could put the desirable post code in this scratch pad [16:0] at reset vector to stop at certain location in BIOS execution. To make forward progress clear the scratch pad.

7.6.3 Error Codes

An error code is a hexadecimal value indicating a fatal error condition, and is a useful tool for debugging. An error code is sent to IO port 80h (or a different IO port defined by input parameter debugPort), and stored into a 32-bit scratchpad register at BIOSNONSTICKYSCRATCHPAD8 in following format:

- [31:24] Major Error Code
- [23:16] Minor Error Code
- [15:8] Major POST Code (see POST Code section)
- [7:0] Minor POST Code (see POST Code section)

The MrcStart() routine returns a 32-bit value in the EAX register. After successful memory initialization, the return value in EAX is zero. A non-zero value in EAX indicates abnormal completion of MRC execution. MrcStart() routine also copies the mrcHost data structure to the buffer using the pointer provided as input parameter by the caller. For serial MRC execution it is stored in SBSP uncore. For parallel MRC execution it is stored in all populated sockets.

A fatal error causes MRC to abort execution hence no memory is successfully initialized and functional. For fatal errors, MRC first writes the major/minor error codes into a scratchpad register as described above. This allows the platform BMC (if available) to access the error codes via SMBus interface. After that, the MRC behavior is determined by the HALT_ON_ERROR directive in the file MemPLATFORM.H. When set to 1 (default), the MRC will enter an infinite loop that displays the major POST code for one second, the minor POST code for half a second, the major error code for half a second and then the minor error code for half a second to the debug port. When set to 0, the MRC will return control to the caller with the major/minor error codes in AH/AL registers of the CPU.

The MRC fatal error handler FatalError() is part of the MemHOOKs.C file which can be customized by OEMs to meet their specific needs.

The fatal error codes are listed below. Please also refer to MemPLATFORM.H file for latest update. Note that some error or warning codes defined may not be used by the current MRC release.

Table 7-2 MRC Fatal Error Code

Fatal Error Nomenclature (As defined in MRC header file)	Major Code	Minor Code	Error Description
ERR_NO_MEMORY	0E8h		
ERR_NO_MEMORY_MINOR_NO_MEMORY		01h	1) No memory was detected via SPD read. No warning log entries available. 2) invalid config that causes no operable memory. Refer to warning log entries for details.
ERR_NO_MEMORY_MINOR_ALL_CH_DISABLED		02h	Memory on all channels of all sockets are disabled due to hardware memtest error
ERR_NO_MEMORY_MINOR_ALL_CH_DISABLED_MIXED		03h	No memory installed. All channels are disabled.
ERR_LT_LOCK	0E9h		Memory locked due to LT "Bricking"
ERR_DDR_INIT	0EAh		DDR3 training did complete successfully
ERR_RD_DQ_DQS		01h	Error on read DQ/DQS init
ERR_RC_EN		02h	Error on Receive Enable
ERR_WR_LEVEL		03h	Error on Write Leveling
ERR_WR_DQ_DQS		04h	Error on write DQ/DQS
ERR_MEM_TEST	0EBh		Memory test failure
ERR_MEM_TEST_MINOR_SOFTWARE		01h	Software memtest failure.
ERR_MEM_TEST_MINOR_HARDWARE		02h	Hardware memtest failed .
ERR_MEM_TEST_MINOR_LOCKSTEP_MODE		03h	Hardware Memtest failure in Lockstep Channel mode requiring a channel to be disabled. This is a fatal error which requires a reset and calling MRC with a different RAS mode to retry.
ERR_VENDOR_SPECIFIC	0ECh		RESERVED for future use
ERR_DIMM_COMPAT	0EDh		UDIMMs and RDIMMs are both present DIMM vendor-specific errors



Fatal Error Nomenclature (As defined in MRC header file)	Major Code	Minor Code	Error Description
ERR_MIXED_MEM_TYPE		01h	Different dimm types are detected installed in the system
ERR_INVALID_POP		02h	Violation of population rules
ERR_INVALID_POP_MINOR_QR_AND_3RD_SLOT		03h	The 3rd DIMM slot can not be populated when QR DIMMs are installed
ERR_INVALID_POP_MINOR_UDIMM_AND_3RD_SLOT		04h	UDIMMs and SODIMMs are not supported in the third DIMM slot
ERR_INVALID_POP_MINOR_UNSUPPORTED_VOLTAGE		05h	Unsupported DIMM Voltage
ERR_DDR3_DDR4_MIXED		06h	DDR3 and DDR4 dimms cannot be mixed
ERR_MIXED_SPD_TYPE		07h	Mixing of 256 byte and 512 byte SPD devices is not supported
ERR_MISMATCH_DIMM_TYPE		08h	Memory type mis-match
ERR_MRC_COMPATIBILITY	0EEh		Number of HAs found in system greater than MAX_HA defined in MRC build
ERR_MRC_DIRE_NONECC		01h	RESERVED for future use
ERR_MRC_STRUCT	0EFh		Indicates a CLTT table structure error. A DIMM is populated in the 3rd slot when quad rank DIMM is present in the channel
ERR_INVALID_BOOT_MODE		01h	Boot mode is unknown
ERR_INVALID_SUB_BOOT_MODE		02h	Sub boot mode is unknown
ERR_SET_VDD	0F0h		
ERR_UNKNOWN_VR_MODE		01h	Invalid VR mode, unable to set DRAM VDD
ERR_IOT_MEM_BUFFER	0F1h		ERROR: IOT Memory Buffer has not been allocated memory

In addition to fatal errors, the MRC supports promotion of non-fatal errors based on the PROMOTE_WARNING directive in MRCPLATFORM.H. When set, the MRC treats any warning as a fatal error. When clear, the MRC logs the warning and continues with memory initialization. The complete list of the MRC warning codes is described in the warning log section below.

7.6.4 Warning Log

For non-fatal errors that do not prevent MRC from initialing all or part of the installed memory, MRC will create Warning Log entries in the MRC output structure (refer to 7.7 for details).

One possible use of the mrcHost structure is to extract MRC warning information. The main BIOS can detect and report non-fatal error conditions after user-friendly interfaces become available. Output devices such as the serial port or video controller can be used to avoid end-user confusion associated with the traditional port 80h debug and beep codes.

During memory detection and initialization, the MRC will log non-fatal errors to the **warningLog** structure referenced by mrcHost.status. This structure field contains warning log entries such as incompatible DIMM type or mismatched DIMM pair. A maximum of 64 warning log entries are available. The actual number of warning log entries created by the MRC is given by the index field (a value of zero means no warnings were logged). Each warning log entry contains a pre-defined warning code. The format of the warning data field is specific to the warning code, but most entries log the node, channel, DIMM number that triggered the warning. The generic warningLog structure definition and a complete list of warning codes are provided below.

```
// Warning log
#define MAX_LOG 64

struct warningEntry {
    U32 code;           // Pre-defined Warning Code
    U32 data;           // Format depends on the Warning Code
};

struct warningLog {
    U32 index;          // Number of warningEntry actually logged
    struct warningEntry log[MAX_LOG];
};
```

Table 7-3 MRC Warning Code

Warning (As defined in MRC header file)	Descriptions	Major Code	Minor Code	Data (DWord)			
		31:16	15:0	31:24	23:16	15:8	7:0
WARN_RDIMM_ON_UDIMM	RDIMMLRDIMM is plugged into a UDIMM only board <i>MRC Behavior:</i> <i>Channel is mapped out.</i>	02h		NODE	CH	DIMM	X
WARN_UDIMM_ON_RDIMM	UDIMM is plugged into a RDIMMLRDIMM only board <i>MRC Behavior:</i> <i>Channel is mapped out.</i>	02h		NODE	CH	DIMM	X
WARN_SODIMM_ON_RDIMM	SODIMM is plugged into a RDIMM/LRDIMM only system <i>MRC Behavior:</i> <i>Channel is mapped out.</i>	03h		NODE	CH	DIMM	X
WARN_4Gb_FUSE	Support for 4Gb devices has been fused off <i>MRC Behavior:</i> <i>Channel is mapped out.</i>	04h		NODE	CH	DIMM	X
WARN_8Gb_FUSE	Support for 8Gb devices has been fused off <i>MRC Behavior:</i> <i>Channel is mapped out.</i>	05h		NODE	CH	DIMM	X
WARN_IMC_DISABLED	IMC was disabled by MRC	06h		NODE	X	X	
WARN_DIMM_COMPAT	DIMM is not compatible with the IMC memory controller. <i>MRC Behavior:</i> <i>Channel is mapped out.</i>	07h		NODE	CH	DIMM	X
WARN_DIMM_COMPAT_MINOR_X16_COMBO	RESERVED for future use		01h	NODE	CH	DIMM	X
WARN_DIMM_COMPAT_MINOR_MAX_RANKS	Max number of ranks exceeded on the channel. MRC has disabled this entire channel. <i>MRC Behavior:</i> <i>Channel is mapped out.</i>		02h	NODE	CH	DIMM	X
WARN_DIMM_COMPAT_MINOR_QR	QR DIMM is not at Slot0 while SR/DR DIMMs are present in the channel. MRC has disabled this entire channel.		03h	NODE	CH	DIMM	X
WARN_DIMM_COMPAT_MINOR_NOT_SUPPORTED	Incompatible DDR3 DIMM module (type/org/tech/speed etc. not supported). MRC has disabled this entire channel.		04h	NODE	CH	DIMM	X
WARN_RANK_NUM	The number of ranks on this device is not supported <i>MRC Behavior:</i> <i>Channel is mapped out after all DIMM's detected.</i>		05h	NODE	CH	DIMM	X
WARN_TOO_SLOW	This DIMM does not support DDR3-1067 or higher <i>MRC Behavior:</i> <i>Channel is mapped out.</i>		06h	NODE	CH	DIMM	X
WARN_DIMM_COMPAT_MINOR_ROW_ADDR_ORDER	LRDIMM A16 usage is not symmetrical on channel <i>MRC Behavior:</i> <i>Channel is mapped out after all DIMM's detected.</i>		07h	NODE	CH	DIMM	X
WARN_CHANNEL_CONFIG_NOT_SUPPORTED	configuration is not supported. Map out this channel. <i>MRC Behavior:</i> <i>Channel is mapped out.</i>		08h	NODE	CH	X	X



Warning (As defined in MRC header file)	Descriptions	Major Code	Minor Code	Data (DWord)			
		31:16	15:0	31:24	23:16	15:8	7:0
WARN_CHANNEL_MIX_ECC_NONE CC	ECC and non ECC dimms were mixed. ECC is disabled for the entire system if one DIMM is found that does not support it. <i>MRC Behavior:</i> <i>Channel will be disabled if ECC Mix was not enabled in the MRC input, if MIX was allowed on Input, then ECC will be disabled.</i>		09h	NODE	CH	DIMM	X
WARN_DIMM_VOLTAGE_NOT_SUPPORTED	DDR4 voltage not supported <i>MRC Behavior:</i> <i>Channel is mapped out.</i>		0Ah	NODE	CH	DIMM	X
WARN_DIMM_COMPAT_TRP_NOT_SUPPORTED	RESERVED for future use		0Bh	NODE	CH	DIMM	X
WARN_LOCKSTEP_DISABLE	Lockstep Channel mode was requested but could not be honored	09h		X	X	X	X
WARN_LOCKSTEP_DISABLE_MINO R_RAS_MODE	Unable to enable Lockstep mode because ECC is disabled. Switch to independent channel mode. (2) <i>MRC Behavior:</i> <i>Warning logged. Mirroring and lockstep will be disabled.</i>		01h	X	X	X	X
WARN_LOCKSTEP_DISABLE_MINO R_MISMATCHED	Mismatched DIMM pairs found accross channels. Switched to independent channel mode. mrchost.RASMode contains the bitmap of supported RAS modes for the current DIMM population. <i>MRC Behavior:</i> <i>Warning logged. Lockstep mode disabled.</i>		02h	X	X	X	X
WARN_LOCKSTEP_DISABLE_MINO R_MEMTEST_FAILED	Memory test failed		03h	X	X	X	X
WARN_USER_DIMM_DISABLE	DIMM was disabled by MRC. See minor code below for specific reasons.	0Ah		NODE	CH	X	X
WARN_USER_DIMM_DISABLE_QUA D_AND_3DPC	3-DIMM-Per-Channel and Quak Rank DIMM were found on the same CPU socket (unsupported config). Channel with Quad Rank DIMM is disabled by MRC.		01h	NODE	CH	X	X
WARN_USER_DIMM_DISABLE_MEM TEST	DIMM was disabled by the MRC as a reslut of previous DIMM in the channel being disabled due to error (the DIMM itself is not necessarily bad)		02h	NODE	CH	DIMM	X
WARN_MEMTEST_DIMM_DISABLE	DIMM was disabled due to MemTest errors. See note (3) below	0Bh		NODE	CH	DIMM	X
WARN_MIRROR_DISABLE	Mirror mode was requested but could not be honored. Memtest failure resulted in a channel being disabled. Switch to independent channel mode. <i>MRC Behavior:</i> <i>Mirroring gets disabled.</i>	0Ch		X	X	X	X
WARN_MIRROR_DISABLE_MINOR_ RAS_DISABLED	Unable to enable Mirror mode because ECC is disabled. Switch to independent channel mode. (2) <i>MRC Behavior:</i> <i>Warning logged. Mirroring gets disabled.</i>		01h	X	X	X	X
WARN_MIRROR_DISABLE_MINOR_ MISMATCH	Mismatched DIMM pairs found accross channels. Switch to independent channel mode. <i>MRC Behavior:</i> <i>Warning logged. Mirroring gets disabled.</i>		02h	X	X	X	X
WARN_MIRROR_DISABLE_MINOR_ MEMTEST	Mirror mode was disabled due to memory test failure		03h	X	X	X	X
WARN_MEM_LIMIT (1)	IMC memory decode limit was reached before all memory could be allocated. (1) <i>MRC Behavior:</i> <i>Warning logged.</i>	0Dh		X	X	X	X



Warning (As defined in MRC header file)	Descriptions	Major Code	Minor Code	Data (DWord)			
		31:16	15:0	31:24	23:16	15:8	7:0
WARN_INTERLEAVE_FAILURE	Interleave mode failure	0Eh					
WARN_SAD_RULES_EXCEEDED	Number of SAD rules exceeds <i>MRC Behavior:</i> <i>Warning logged. All present memory does NOT get mapped.</i>		01h	X	X	X	X
WARN_TAD_RULES_EXCEEDED	Number of TAD rules exceeds <i>MRC Behavior:</i> <i>Warning logged. All present memory does NOT get mapped.</i>		02h	NODE	X	X	X
WARN_RIR_RULES_EXCEEDED	Number of RIR rules exceeds <i>MRC Behavior:</i> <i>Warning logged. All present memory does NOT get mapped.</i>		03h	NODE	CH	X	X
WARN_TAD_OFFSET_NEGATIVE	Negative TAD offset <i>MRC Behavior:</i> <i>Warning logged.</i>		04h	NODE	X	X	X
WARN_TAD_LIMIT_ERROR	TAD Limit > SAD Limit <i>MRC Behavior:</i> <i>Warning logged.</i>		05h	NODE	X	X	X
WARN_INTERLEAVE_3WAY	Interleave 3 Way		06h	NODE	X	X	X
WARN_A7_MODE_AND_3WAY_CH_I NTERLV	RESERVED for future use		07H	NODE	X	X	X
WARN_SPARE_DISABLE	<ul style="list-style-type: none"> Unable to enable Spare mode because ECC is disabled. Switch to independent channel mode. (2) Mismatched DIMM pairs found accross channels. Switch to independent channel mode. <i>MRC Behavior:</i> <i>Warning logged. Sparing gets disabled.</i>	10h		X	X	X	X
WARN_PTRLSR_DISABLE	Patrol Scrub was disabled	11h		X	X	X	X
WARN_UNUSED_MEMORY	Unused memory is populated on channel 2 in Lockstep or Mirroring mode.	12h		NODE	CH	X	X
WARN_UNUSED_MEMORY_MIRRO R	Unused memory is populated on channel 2 in mirror mode		01h	NODE	2	X	X
WARN_UNUSED_MEMORY_LOCKS TEP	Unused memory is populated on channel 2 in Lockstep mode		02h	NODE	2	X	X
WARN_RD_DQ_DQS	A Read DQ/DQS failure has occurred during training. The failing Channel was disabled	13h		NODE	CH	DIMM	X
WARN_RD_RCVEN	A tRLCoarse failure has occurred during DDR training. The failing Channel was disabled. <i>MRC Behavior:</i> <i>Channel is disabled.</i>	14h		NODE	CH	X	X
WARN_ROUNDTRIP_EXCEEDED	Round Trip delay of %d exceeds the limit of %d <i>MRC Behavior:</i> <i>Channel is disabled.</i>		01h	NODE	CH	DIMM	RANK
WARN_WR_LEVEL	A write leveling failure has occurred during training. (1)	15h		NODE	CH	DIMM	X
WARN_WR_FLYBY_CORR	Fault Parts Tracking write Fly-by correctable error		00h	NODE	CH	X	X
WARN_WR_FLYBY_UNCORR	Fault Parts Tracking write Fly-by uncorrectable error		01h	NODE	CH	X	X
WARN_WR_FLYBY_DELAY	RESERVED for future use		02h	NODE	CH	X	X
WARN_WR_DQ_DQS (1)	A write DQ/DQS failure has occurred during training. (1) <i>MRC Behavior:</i> <i>Fatal error. Channel is mapped out.</i>	16h		NODE	CH	DIMM	X
WARN_DIMM_POP_RULE	Improper DIMM population	17h		NODE	CH	DIMM	X



Warning (As defined in MRC header file)	Descriptions	Major Code 31:16	Minor Code 15:0	Data (DWord)			
				31:24	23: 16	15:8	7:0
WARN_DIMM_POP_RULE_MINOR_OUT_OF_ORDER	DIMM is populated out of order and that it will not be used. If slot 0 is empty then the channel gets disabled, if slot 1 is empty but slot 0 and slot 2 are populated then MRC will try to boot with the DIMM in slot 0 while ignoring the DIMM in slot 2.		01h	NODE	CH	DIMM	X
WARN_DIMM_POP_RULE_MINOR_INDEPENDENT_MODE	Lockstep/Mirror mode not enabled due to unused DIMM on Channel 2, and MRC input RAS_TO_INDP_EN = 1. Switch to Independent Channel mode.		02h	NODE	2	X	X
WARN_CLTT_DISABLE	CLTT was requested but could not be honored	18h					
WARN_CLTT_MINOR_NO_TEMP_SENSOR	A DIMM without Temp Sensor was found		01h	NODE	CH	DIMM	X
WARN_CLTT_MINOR_CIRCUIT_TEST_FAILED	A DIMM failed Temp Sensor circuit test		02h	NODE	CH	DIMM	X
WARN_THROT_INSUFFICIENT	Indicates throttling is not sufficient for this DIMM due to MRC calculation.	19h		NODE	CH	DIMM	X
WARN_CLTT_DIMM_UNKNOWN	A DIMM of an unknown category is found when looking up a pre-defined category table (DIMM type, rawcard, heat spreader, planner, etc). Use a default category (category 11 or 27 depending on DIMM type)	1Ah		NODE	CH	DIMM	X
WARN_DQS_TEST	DQS training failure encountered	1Bh		X	X	X	X
WARN_MEM_TEST	Hardware Memtest failed and the DIMM is disabled	1Ch		NODE	CH	DIMM	X
WARN_CLOSED_PAGE_OVERRIDE	Page override	1Dh		NODE	X	X	X
WARN_DIMM_VREF_NOT_PRESENT	DIMM Vref controller circuit (DCP) not detected	1Eh		NODE	X	X	X
WARN_LV_STD_DIMM_MIX	Low voltage DDR3 problem encountered.	20h		NODE	X	X	X
WARN_LV_2QR_DIMM	2 LV QR detected	21h		NODE	X	X	X
WARN_LV_3DPC	LV 3 DPC detected	22h		NODE	X	X	X
WARN_CMD_ADDR_PARIT_ERR	RESERVED for future use	23h		NODE	X	X	X
WARN_COD_HA_NOT_ACTIVE	COD Enabled but memory not behind each HA	25h		NODE	X	X	X
WARN_CMD_CLK_TRAINING	Eye width is too small. Channel is disabled.	26h		NODE	CH	X	X
WARN_FPT_CORRECTABLE_ERROR	FTP correctable error <i>MRC Behavior: mentioned in the logfile.</i>	30h					
WARN_FPT_UNCORRECTABLE_ERROR	FTP uncorrectable error <i>MRC Behavior: Rank is disabled.</i>	31h					
WARN_FPT_MINOR_WR_FLYBY	FTP: failed Write FlyBy		00h	NODE	CH	DIMM	RANK
WARN_FPT_MINOR_RD_DQ_DQS	FTP: failed Read DqDqs		13h	NODE	CH	DIMM	RANK
WARN_FPT_MINOR_RD_RCVEN	Receive Enable training failure		14h	NODE	CH	DIMM	RANK
WARN_FPT_MINOR_WR_LEVEL	FTP failed Write Levelling		15h	NODE	CH	DIMM	RANK
WARN_FPT_MINOR_WR_DQ_DQS	FTP: failed Write DqDqs		16h	NODE	CH	DIMM	RANK
WARN_FPT_MINOR_DQS_TEST	FTP: failed DQS Test		1Bh	NODE	CH	DIMM	RANK
WARN_FPT_MINOR_MEM_TEST	FTP minor correctable memtest		1Ch	NODE	CH	DIMM	RANK
WARN_FPT_MINOR_LRDIMM_TRAINING	FTP minor LRDIMM training		24h	NODE	CH	DIMM	RANK
WARN_FPT_MINOR_VREF_TRAINING	Unable to find the txvref for the eye. Warning for both correctable and uncorrectable errors. Channel is disabled.		26h	NODE	CH	DIMM	RANK
WARN_NO_SUFFICIENT_SPARERANK	<i>MRC Behavior: Warning is logged.</i>	32h		NODE	CH	X	X
WARN_MEM_CONFIG_CHANGED	Timing overrides are enabled but the DIMM configuration has changed. Memory overrides will be disabled	40h		X	X	X	X
WARN_MEM_OVERRIDE_DISABLED	If MEM_OVERRIDE_EN is enabled but the DIMM configuration has changed, this warning indicates that the MRC has disabled memory overrides.		01h	X	X	X	X
WARN_MCA_UNCORRECTABLE_ERROR	MC uncorrectable error caused machine check	50h		X	X	X	X
WARN_DM_TEST_ERROR_CODE	Major DMT error	60h		X	X	X	X



Warning (As defined in MRC header file)	Descriptions	Major Code	Minor Code	Data (DWord)			
		31:16	15:0	31:24	23:16	15:8	7:0
WARN_DM_TEST_PARSE_ERROR_MINOR_CODE	Parse error		01h	X	X	X	X
WARN_DM_TEST_CONFIGURATION_ERROR_MINOR_CODE	Mix Dimm configuration unsupported on DRAM maintenance test		02h	X	X	X	X
WARN_DM_TEST_EXECUTION_ERROR_MINOR_CODE	Couldn't calculate number loopcounts for subsequences		03h	X	X	X	X

Notes:

- (1) Currently unused code as a placeholder. May be used or removed in the future.
- (2) ECC gets disabled in the memory controller by MRC due to any of the following reasons:
 - MRC input option ECC_CHECK_EN = 0
 - The processor sku does not support ECC
 - One or more DIMMs detected does not support ECC
- (3) A DIMM that fails MemTest or DDR training will be disabled. The failing DIMM and all DIMMs behind it on the same channel will also be disabled, or mapped out (their DOD's cleared, the Rank Present bits cleared and they will not have memory address region mapped to them). MemTest is run again after the DIMMs are mapped out. If a RAS mode(lockstep, mirroring, or sparing) is requested then MRC will attempt to retain the RAS mode, unless DIMM0 fails, in which this case the MRC will revert to independent channel mode. For lockstep mode, however, switching to independent channel mode in this case is impossible due to the need of a reset. When a DIMM is disabled, a warning log entry is created with
- (4) If necessary, DIMM operating frequency can be "clipped" via BIOS Setup Option

7.7 MRC Output Results to Physical Memory

The MRC initializes a large number of data structures in CPU NEM cache (CAR) during memory detection and initialization. These data structures contain a wealth of information accumulated during MRC execution and can be copied to physical memory for main BIOS consumption after MRC completes.

It is important to preserve MRC results in physical memory because the main BIOS will eventually disable the NEM environment and continue execution from memory. Consequently, all the MRC information accumulated in CPU NEM cache (CAR) will be lost. Saving these structures to memory helps the main BIOS report memory configuration and implement advanced error handling features.

7.7.1 Output Structure

The MRC will output data in the SYSHOST structure in the section referenced as memVar. The format of the output data is defined by the **memVar** structure in MEMHOST.H. This structure definition may expand in future MRC releases.

To facilitate relocation of output data, the mrcHost structure does not contain any pointers. Other internal data structures have been nested within the mrcHost structure for a single point of reference. Information about DDR channel topology, DIMM population, and host configuration can be extracted from this data structure. A complete description of the memVar structure can be found in the header file MEMHOST.H. Individual BIOS developers should review the MRC usage of internal data structures to understand the possibilities for main BIOS usage.



8 Compatible Data Structure

8.1 Requirements

Intel BIOS reference code implements system validation features that produce significant amounts of data. Customers and suppliers need a compatible method to access and parse this data with generic tools. The access method must support native access from applications running on the target system as well as remote access via ITP (or serial connection) from applications running on a host system. The data should be accessible from the time it is produced throughout the BIOS boot flow and into the OS. To maintain compatibility with future platforms, a pointer to the data structure must be provided via standard mechanisms defined in the Advanced Configuration and Power Interface (ACPI) specification.

The data structure format must be specified and associated with a unique version number such that non-BIOS applications can discover the format and maintain backward compatibility with old revisions. Forward compatibility is not a requirement, but a secondary version number can denote when fields have been appended to the end of the compatible structure. The data structure must support a reliable mechanism to verify data integrity, such as a Cyclic Redundancy Check (CRC) algorithm.

8.2 Proposed Solution

Intel BIOS reference code shall define a compatible data structure using the C programming language and compiler settings for Intel® IA32 Architecture. The BIOS data structure shall consist of three sections: a compatible header, a versioned data range, and an optional OEM data range. The compatible header shall contain the following information: an 8-byte signature string, a total structure size, a 16-bit CRC covering the structure size, primary and secondary versions, and an optional OEM offset.

The versions shall uniquely define the format of the data range such that an application can cast the memory range with the associated C structure and decode the data fields. The secondary version number shall be incremented when data fields are appended to the previous version of the data structure. The versions do not apply to the OEM data range.

BIOS reference code shall reserve a memory range on the System BSP stack while operating from cache in No-Eviction Mode (NEM). The memory range shall be large enough to store the compatible data structure. The BIOS reference code shall initialize fields in the data structure, update the CRC value, and store a pointer to the data structure in a scratchpad register in the processor. To determine BIOS execution context, a remote host application can optionally read the current BIOS checkpoint from a separate scratchpad register. The meaning of checkpoints on a given platform shall be consistent within BIOS reference code, but checkpoints may vary across different System BIOS implementations and different platforms.

After BIOS reference code completes, the System BIOS shall copy the data structure from the NEM stack to an intermediate system memory location and update the pointer in the scratchpad register. The data structure on the NEM stack will be invalidated by tearing down the NEM stack. The System BIOS shall reserve the intermediate memory range using the POST Memory Manager, EFI memory management services, and internal mechanisms so that subsequent steps during BIOS POST do not corrupt the data structure. If the System BIOS needs to update fields within the data structure, it shall be responsible for recalculating the CRC after the updates. The data structure must be relocated to different memory addresses, so pointers to fields within the data structure should be avoided. Instead, offsets can be used relative to the base address of the data structure. External pointers should also be defined as offsets of type UINT32 or UINT64 (relative to base address 0).

When the system BIOS establishes the final system memory map and location of the ACPI memory regions, the data structure shall be copied from the intermediate memory location to its final resting place in AddressRangeReserved, an ACPI Type 2 memory region below 4 GB. The system BIOS must reserve a memory region that is large enough to accommodate the size of the BIOS data structure rounded up to the next 4 KB page size and aligned to a 4 KB address boundary. The system BIOS shall report the physical address range of the Type 2 memory region as required by the ACPI specification. This includes software Interrupt 15h - function AX = E820h, or EFI GetMemoryMap if applicable.

The system BIOS shall initialize a custom ACPI table containing a pointer to the physical base address of the BIOS data structure within the Type 2 memory region. The system BIOS shall then update the scratchpad register with the physical address of the BIOS data structure and free the intermediate memory location. All subsequent accesses to the BIOS data structure should be directed to the final runtime location in the ACPI Type 2 region.

8.3 Compatible ACPI Table

The pointer to the BIOS data structure shall be defined in a custom ACPI table to provide compatibility across multiple platforms. The Root System Description Table (RSDT) shall reference a custom OEM table identified by the unique signature "BDAT". The BDAT table shall conform to the standard ACPI header and contain a Global Address Structure that defines the 64-bit physical base address of the BIOS data structure. An OS driver may be required to access the custom ACPI table and to load pages containing the BIOS data structure.

The following C code is a sample implementation of the BDAT table based on the EFI Developer Kit. The BdatGas field and the table checksum must be updated at boot time based on the address of the BIOS data structure.



```
#pragma pack(1)

typedef unsigned char      UINT8;
typedef unsigned short    UINT16;
typedef unsigned long     UINT32;
typedef unsigned long long UINT64;

#define EFI_SIGNATURE_16(A, B)      ((A) | (B << 8))
#define EFI_SIGNATURE_32(A, B, C, D) (EFI_SIGNATURE_16 (A, B) | (EFI_SIGNATURE_16 (C, D) << 16))

//
// Common ACPI description table header.  This structure prefaces most ACPI tables.
//
typedef struct {
    UINT32  Signature;
    UINT32  Length;
    UINT8   Revision;
    UINT8   Checksum;
    UINT8   OemId[6];
    UINT64  OemTableId;
    UINT32  OemRevision;
    UINT32  CreatorId;
    UINT32  CreatorRevision;
} EFI_ACPI_DESCRIPTION_HEADER;

//
// ACPI 3.0 Generic Address Space definition
//
typedef struct {
    UINT8   AddressSpaceId;
    UINT8   RegisterBitWidth;
    UINT8   RegisterBitOffset;
    UINT8   AccessSize;
    UINT64  Address;
} EFI_ACPI_3_0_GENERIC_ADDRESS_STRUCTURE;

//
// BIOS Data ACPI structure
//
typedef struct {

    EFI_ACPI_DESCRIPTION_HEADER    Header;
    EFI_ACPI_3_0_GENERIC_ADDRESS_STRUCTURE  BdatGas;

} EFI_BDAT_ACPI_DESCRIPTION_TABLE;

//
// BIOS Data Parameter Region Generic Address
// Information
//
#define EFI_BDAT_ACPI_POINTER      0x0

//
// BIOS Data Table
//
EFI_BDAT_ACPI_DESCRIPTION_TABLE  BiosDataTable = {
    EFI_SIGNATURE_32('B','D','A','T'), // Signature
    sizeof (EFI_BDAT_ACPI_DESCRIPTION_TABLE), // Length
    0x01, // Revision [01]
    //
    // Checksum will be updated during boot
    //
    0, // Checksum
    ' ', // OEM ID
    ' ',
    ' ',
    ' ',
    ' ',
    ' ',
    ' ',
    ' ',
    0, // OEM Table ID
}
```



```

0, // OEM Revision [0x00000000]
0, // Creator ID
0, // Creator Revision
0, // System Memory Address Space ID
0,
0,
0,
0,

// Pointer will be updated during boot
EFI_BDAT_ACPI_POINTER,
};

#pragma pack()

```

8.4 Scratchpad

A scratchpad configuration register shall be reserved to facilitate usage of the BIOS data structure during boot time and by remote host applications. The scratchpad register will contain a 32-bit physical address pointer to the BIOS data structure and must be qualified by the current BIOS checkpoint. The BIOS checkpoint will be located in a different scratchpad register. A NULL pointer of 0 is not valid and must not be de-referenced.

The scratchpad registers will be located in the System BSP Uncore domain as defined in Section 4.5. BDAT feature uses BIOSNonStickyScratchpad5 and BIOSNonStickyScratchpad7.

8.5 Compatible BIOS Structure Header

The BIOS data structure shall begin with a compatible header so that an application can determine the remaining structure format and check the data integrity. The header shall contain the following information: an 8-byte signature string, a total structure size, a 16-bit CRC, primary and secondary versions, and an optional OEM offset.

The signature string shall be initialized to the ASCII sequence "BDATHEAD". The CRC shall be calculated over the specified size of the BIOS data structure, assuming that the CRC field itself contains a value of 0. The 16-bit CRC algorithm shall be compatible with the JEDEC DDR3 Serial Presence Detect (SPD) specification for bytes 126 – 127.

The primary and secondary versions shall uniquely define the format of the data range such that an application can cast the memory range with the associated C structure and decode the data fields. The secondary version number shall be incremented when data fields are appended to the previous version of the data structure. The secondary version number can be recycled when the primary version changes.

The OEM offset provides an optional mechanism for OEMs to customize the BIOS data structure without affecting compatibility of the versioned data range. The version numbers do not apply to the OEM data range, although fields in the versioned data range can be initialized by the OEM system BIOS. The OEM offset is provided mainly as a courtesy for customers that wish to use the BIOS data structure mechanism to transfer information to an OS driver. The format of the OEM data range is outside the scope of this specification.

The header may not be aligned during BIOS use but will be aligned to a 4 KB page boundary when relocated in Type 2 memory for OS use. The header format shall not change across different platform generations. The following data structure defines the compatible header.

```

#pragma pack(1)

typedef struct {
    UINT8    BiosDataSignature[8]; // "BDATHEAD"
    UINT32    BiosDataStructSize; // sizeof BDAT_STRUCTURE
    UINT16    Crc16; // 16-bit CRC of BDAT_STRUCTURE (calculated with 0 in thisfield)
    UINT16    Reserved
    UINT16    PrimaryVersion; // Primary version
    UINT16    SecondaryVersion; // Secondary version
    UINT32    OemOffset; // Optional offset to OEM-defined structure

    UINT32    Reserved1;
    UINT32    Reserved2;
} BDAT_HEADER_STRUCTURE;

#pragma pack()

```

8.6 BIOS Structure Definitions

The definition of the compatible BIOS data structure shall be managed by the BIOS reference code team. Any changes to the data structure identified by the primary and secondary version numbers must be approved by the BIOS reference code team and the external tools team. The version numbers can be combined into a shorthand notation of “primary.secondary” version when referencing the BIOS data structure. The following code defines a preliminary version 2.0 data structure, BDAT_SYSTEM_STRUCTURE. *Please refer to BDAT.H in the Reference Code for latest structure definition.*

Version 2.0: The following code defines version 2.0 of the BDAT_SYSTEM_STRUCTURE.

```
#pragma pack(1)

#define MC_MAX_NODE          8           // Max memory nodes per system
#define MAX_CH               4           // Max channels per memory node
#define MAX_DIMM             3           // Max DIMM per channel
#define MAX_RANK_DIMM        4           // Max ranks per DIMM
#define MAX_STROBE           18          // Number of strobe groups
#define MAX_SPD_BYTE         256         // Number of bytes in Serial EEPROM

typedef struct {
    UINT8  rxDqLeft;           // Units = 1/64 QCLK
    UINT8  rxDqRight;
    UINT8  txDqLeft;
    UINT8  txDqRight;
    UINT8  cmdLeft;
    UINT8  cmdRight;
    UINT8  rxVrefLow;
    UINT8  rxVrefHigh;
    UINT8  txVrefLow;
    UINT8  txVrefHigh;
} BDAT_RANK_MARGIN_STRUCTURE;

typedef struct {
    UINT16  recEnDelay[MAX_STROBE];      // Units = 1/64 QCLK
    UINT16  wldDelay[MAX_STROBE];
    UINT8   rxDqDelay[MAX_STROBE];
    UINT8   txDqDelay[MAX_STROBE];
    UINT8   clkDelay;
    UINT8   ctlDelay;
    UINT8   cmdDelay[3];
    UINT8   ioLatency;
    UINT8   roundtrip;
} BDAT_RANK_TRAINING_STRUCTURE;

typedef struct {
    UINT16  mr0;
    UINT16  mr1;
    UINT16  mr2;
    UINT16  mr3;
} BDAT_RANK_MRS_STRUCTURE;

typedef struct {
    UINT8   rankEnabled;
    BDAT_RANK_MARGIN_STRUCTURE  rankMargin;
    BDAT_RANK_TRAINING_STRUCTURE  rankTraining;
    BDAT_RANK_MRS_STRUCTURE      rankMRS;
} BDAT_RANK_STRUCTURE;

typedef struct {
    UINT8   valid[MAX_SPD_BYTE/8];      // Each valid bit maps to SPD byte
    UINT8   spdData[MAX_SPD_BYTE];
} BDAT_SPD_STRUCTURE;

typedef struct {
    UINT8   dimmEnabled;
    BDAT_RANK_STRUCTURE      rankList[MAX_RANK_DIMM];
    BDAT_SPD_STRUCTURE        spdBytes;
} BDAT_DIMM_STRUCTURE;

typedef struct {
```




```

        UINT8                chEnabled;
        UINT8                numDimmSlot;
        BDAT_DIMM_STRUCTURE  dimmList[MAX_DIMM];
    } BDAT_CHANNEL_STRUCTURE;

typedef struct {
    UINT8                imcEnabled;
    UINT16               ddrFreq;           // Units of MT/s
    UINT16               ddrVoltage;       // Units of mV
    UINT16               imcDid;
    UINT8                imcRid;
    BDAT_CHANNEL_STRUCTURE channelList[MAX_CH];
} BDAT_NODE_STRUCTURE;

typedef struct {
    UINT32               refCodeRevision;
    UINT32               marginLoopCount;   // Units of cache line
    BDAT_NODE_STRUCTURE nodeList[MC_MAX_NODE];
} BDAT_SYSTEM_STRUCTURE;

typedef struct bdatStruct {
    BDAT_HEADER_STRUCTURE bdatHeader;
    BDAT_SYSTEM_STRUCTURE bdatSys;
} BDAT_STRUCTURE;

#pragma pack()

```

Version 2.0: The following pseudo-code defines version 2.0 of the BDAT_SYSTEM_STRUCTURE.

```

#pragma pack(1)

typedef struct {
    UINT8  rxDqLeft;       // Units = piStep
    UINT8  rxDqRight;
    UINT8  txDqLeft;
    UINT8  txDqRight;
    UINT8  rxVrefLow;      // Units = rxVrefStep
    UINT8  rxVrefHigh;
    UINT8  txVrefLow;      // Units = txVrefStep
    UINT8  txVrefHigh;
} BDAT_DQ_MARGIN_STRUCTURE;

typedef struct {
    UINT8  rxDqLeft;       // Units = piStep
    UINT8  rxDqRight;
    UINT8  txDqLeft;
    UINT8  txDqRight;
    UINT8  cmdLeft;
    UINT8  cmdRight;
    UINT8  rxVrefLow;      // Units = rxVrefStep
    UINT8  rxVrefHigh;
    UINT8  txVrefLow;      // Units = txVrefStep
    UINT8  txVrefHigh;
} BDAT_RANK_MARGIN_STRUCTURE;

typedef struct {
    UINT16  recEnDelay[maxStrobe];          // Array of nibble training results per rank
    UINT16  wlDelay[maxStrobe];
    UINT8   rxDqDelay[maxStrobe];
    UINT8   txDqDelay[maxStrobe];
    UINT8   clkDelay;
    UINT8   ctlDelay;
    UINT8   cmdDelay[3];
    UINT8   ioLatency;
    UINT8   roundtrip;
} BDAT_RANK_TRAINING_STRUCTURE;

typedef struct {

```




```
    UINT16 mr0;                // MR0 settings
    UINT16 mr1;                // MR1 settings
    UINT16 mr2;                // MR2 settings
    UINT16 mr3;                // MR3 settings
} BDAT_RANK_MRS_STRUCTURE;

typedef struct {
    UINT8                rankEnabled;        // 0 = Rank disabled
    UINT8                rankMarginEnabled;  // 0 = Rank margin disabled
    UINT8                dqMarginEnabled;    // 0 = Dq margin disabled
    BDAT_RANK_MARGIN_STRUCTURE rankMargin;   // Rank margin data
    BDAT_DQ_MARGIN_STRUCTURE dqMargin[maxDq]; // Array of Dq margin data per rank
    BDAT_RANK_TRAINING_STRUCTURE rankTraining; // Rank training settings
    BDAT_RANK_MRS_STRUCTURE rankMRS;        // Rank MRS settings
} BDAT_RANK_STRUCTURE;

#define MAX_SPD_BYTE        256    // Number of bytes in Serial EEPROM

typedef struct {
    UINT8                valid[MAX_SPD_BYTE/8]; // Each valid bit maps to SPD byte
    UINT8                spdData[MAX_SPD_BYTE]; // Array of raw SPD data bytes
} BDAT_SPD_STRUCTURE;

typedef struct {
    UINT8                dimmEnabled;        // 0 = DIMM disabled
    BDAT_RANK_STRUCTURE rankList[maxRankDimm]; // Array of ranks per DIMM
    BDAT_SPD_STRUCTURE  spdBytes;           // SPD data per DIMM
} BDAT_DIMM_STRUCTURE;

typedef struct {
    UINT8                chEnabled;          // 0 = Channel disabled
    UINT8                numDimmSlot;        // Number of slots per channel on the board
    BDAT_DIMM_STRUCTURE dimmList[maxDimm];  // Array of DIMMs per channel
} BDAT_CHANNEL_STRUCTURE;

typedef struct {
    UINT8                imcEnabled;        // 0 = MC disabled
    UINT16               imcDid;            // MC device Id
    UINT8                imcRid;            // MC revision Id
    UINT16               ddrFreq;           // DDR frequency in units of MHz / 10
                                         // e.g. ddrFreq = 1333 for tCK = 1.5 ns
    UINT16               ddrVoltage;        // Vdd in units of mV
                                         // e.g. ddrVoltage = 1350 for Vdd = 1.35 V
    UINT8                piStep;            // Step unit = piStep * tCK / 2048
                                         // e.g. piStep = 16 for step = 11.7 ps (1/128 tCK)
    UINT16               rxVrefStep;        // Step unit = rxVrefStep * Vdd / 100
                                         // e.g. rxVrefStep = 520 for step = 7.02 mV
    UINT16               txVrefStep;        // Step unit = txVrefStep * Vdd / 100
    BDAT_CHANNEL_STRUCTURE channelList[maxCh]; // Array of channels per socket
} BDAT_SOCKET_STRUCTURE;

typedef struct {
    UINT32               refCodeRevision;   // Matches scratchpad definition
    UINT8                maxNode;           // Max processors per system, e.g. 4
    UINT8                maxCh;             // Max channels per socket, e.g. 4
    UINT8                maxDimm;          // Max DIMM per channel, e.g. 3
    UINT8                maxRankDimm;      // Max ranks per DIMM, e.g. 4
    UINT8                maxStrobe;        // Number of Dqs used by the rank, e.g. 18
    UINT8                maxDq;            // Number of Dq bits used by the rank, e.g. 72
    UINT32               marginLoopCount;   // Units of cache line
    BDAT_SOCKET_STRUCTURE socketList[maxNode]; // Array of sockets per system
} BDAT_SYSTEM_STRUCTURE;

typedef struct bdatStruct {
    BDAT_HEADER_STRUCTURE bdatHeader;
    BDAT_SYSTEM_STRUCTURE bdatSys;
} BDAT_STRUCTURE;

#pragma pack()
```



9 Faulty Parts Tracking

FPT feature is part of the MRC training flow for the following training steps:

- Receive Enable
- Read DQ/DQS
- Write Levelling
- Write DQ/DQS

During these training steps, FPT keeps track of errors detected on a per nibble (x4 width) basis. The results are stored in an array of 1 byte information per Node/Channel/DIMM/Rank/Strobe. The array is defined as:

```
UINT8          faultyParts[MAX_STROBE];
```

And is part of structure:

```
struct ddrRank {};
```

Initially all values in the array are set to 0x00. This means no error was found on any nibble. If a failing nibble is found during any of the training steps listed above it would be marked bitwise as described below:

```
#define FPT_REC_ENABLE_FAILED      (0x01)
#define FPT_RD_DQ_DQS_FAILED      (0x02)
#define FPT_WR_LEVELING_FAILED    (0x04)
#define FPT_WR_DQ_DQS_FAILED      (0x10)
```

If a nibble fails in a particular training step all subsequent training steps will be skipped for that particular nibble. After each training step where a faulty nibble has been found, an evaluation function is called to assess if the faulty nibble(s) create an uncorrectable error. If the result of the evaluation indicates that an uncorrectable error has been detected MRC execution will be stopped, otherwise the MRC will continue to execute as usual.

After MRC execution is complete, the above mentioned array can be read to obtain information on any failing nibbles detected.

This information is also part of the serial debug output. An example serial output with FPT information is shown below:

```
N0.C0.D0.R0.S17: Center Point = 39
N0.C0.D0.R0.S17: WrLevel Logic Delay = 0, Pi = 3
STOP_DATA_WR_LVL_BASIC
END - Wr Lvl Training Time      712 ms
Entering no zone 11
START - Tx Dq Training
Checkpoint Code: Socket 0, 0xB7, 0x04, 0x0000
FAULTY_PARTS_TRACKING: EvaluateFaultyParts status: 1
A warning has been logged! Warning Code = 0x16, Minor Warning Code = 0x0, Data = 0x0
Socket = 0 Channel = 0 DIMM = 0 Rank = 0
FAULTY_PARTS_TRACKING: EvaluateFaultyParts status: 2
N0.C0.D0.R0: WrDq/Dqs Failure! (status = 1)
A warning has been logged! Warning Code = 0x16, Minor Warning Code = 0x0, Data = 0x0
Socket = 0 Channel = 0 DIMM = 0 Rank = 0
```



10 Storage Feature MRC Users Guide

The MRC contains support for 2 different ADR usage models that take two very different paths through the MRC. This guide provides a brief overview of each usage model and describes how the MRC should be configured for each usage model.

ADR+BBU: In this usage model the DIMMs are placed in self-refresh upon system power loss by ADR. Following self-refresh entry power is removed from the entire system except for the DIMMs, which are powered via battery. The system remains in this state until system power is restored. Upon entering the MRC, the DIMMs contain data that must be preserved. Therefore the MRC takes a path very similar to S3Resume to insure that data is not lost.

ADR+NVDIMM: In this usage model the DIMMs are placed in self-refresh upon system power loss by ADR. However, following self-refresh entry, an NVDIMM will isolate itself from the host DDR bus. After isolation, on-DIMM controller will copy all data from DRAM to Flash that is also located on the NVDIMM. During this process, the NVDIMM is powered by a Supercap pack attached to NVDIMM via cable, and the rest of the system is free to completely power down. When power is restored the system will power back up, and the MRC will instruct the NVDIMMs to copy data back from Flash to DRAM.

10.1 MRC Configuration Options

Compile time Flags:

#define	Description
MEM_NVDIMM_EN	Controls if NVDIMM code is enabled or not
NVMEM_FEATURE_EN	Controls if the Enhanced interleave code required by both NVDIMM and ADR is enabled

Input Parameters:

Variable	Description
batterybacked	Controls if a given channel is battery backed or not. Only applies to ADR+BBU Must be 0 for all channels for NVDIMM Usage models
ADREn	Controls if the ADR trigger will be enabled. Should be set for both ADR+BBU and ADR+NVDIMM
check_pm_sts	Controls if the PCH_PM_STS status register is read to determine if the system should recover from ADR Only applies to ADR+BBU Must be 0 for all channels for NVDIMM Usage models
check_platform_detect	Controls if the platform specific PlatformDetectADR function is called determine if the system should recover from ADR Only applies to ADR+BBU Must be 0 for all channels for NVDIMM Usage models
eraseArmNVDIMMS	Controls if the MRC will Erase and Arm all NVDIMM in the system or if it will leave those operations to the OS

Recommended Configuration values

Parameter	ADR+BBU	ADR + NVDIMM
MEM_NVDIMM_EN	Don't Care	Must be Defined
NVMEM_FEATURE_EN	Must be Defined	Must be Defined
batterybacked	Set for battery backed channels	Clear for all channels
ADREn	Must be Set	Must be Set
check_pm_sts	Platform Dependent	Must be Clear
check_platform_detect	Platform Dependent	Must be Clear
eraseArmNVDIMMS	Don't Care	Platform Dependent

10.2 Output Parameters

Parameter	Description																		
nvDimmType	A per DIMM field that indicates if a given DIMM is a normal DIMM or an NVDIMM. A value of 0 indicates that it is either not an NVDIMM or not an NVDIMM recognized by the system. Any other value indicates that it is an NVDIMM.																		
nvDimmStatus	<div><div>A per DIMM field that provides status about the operations performed on the NVDIMM by the MRC. The following bits are defined</div><table><tr><th>Name</th><th>Bit</th><th>Description</th></tr><tr><td>STATUS_RESTORE_SUCCESSFUL</td><td>2</td><td>Indicates that a restore was successfully completed by this NVDIMM.</td></tr><tr><td>STATUS_ARMED</td><td>3</td><td>Indicates that this NVDIMM was successfully armed, and is capable of starting a SAVE operation</td></tr><tr><td>ERROR_DETECT</td><td>8</td><td>Indicates that the MRC was unable to determine the status of the NVDIMM, and did not attempt any further operations on it</td></tr><tr><td>ERROR_RESTORE</td><td>9</td><td>Indicates that the MRC attempted to complete a RESTORE operation, but failed</td></tr><tr><td>ERROR_ARM</td><td>10</td><td>Indicates that the MRC attempted to complete a ARM operation, but failed</td></tr></table></div>	Name	Bit	Description	STATUS_RESTORE_SUCCESSFUL	2	Indicates that a restore was successfully completed by this NVDIMM.	STATUS_ARMED	3	Indicates that this NVDIMM was successfully armed, and is capable of starting a SAVE operation	ERROR_DETECT	8	Indicates that the MRC was unable to determine the status of the NVDIMM, and did not attempt any further operations on it	ERROR_RESTORE	9	Indicates that the MRC attempted to complete a RESTORE operation, but failed	ERROR_ARM	10	Indicates that the MRC attempted to complete a ARM operation, but failed
Name	Bit	Description																	
STATUS_RESTORE_SUCCESSFUL	2	Indicates that a restore was successfully completed by this NVDIMM.																	
STATUS_ARMED	3	Indicates that this NVDIMM was successfully armed, and is capable of starting a SAVE operation																	
ERROR_DETECT	8	Indicates that the MRC was unable to determine the status of the NVDIMM, and did not attempt any further operations on it																	
ERROR_RESTORE	9	Indicates that the MRC attempted to complete a RESTORE operation, but failed																	
ERROR_ARM	10	Indicates that the MRC attempted to complete a ARM operation, but failed																	
NVmemSize	A variable present at the dimm, channel, home agent, socket, or full system levels to determine the capacity of non-volatile memory(from ADR+BBU or ADR+NVDIMM) present at a given level of the hierarchy. These variables in conjunction with the NUMA_AWARE flag can be used by the BIOS to determine the memory map of the system.																		

10.3 Memory map layout

If the NVMEM_FEATURE_EN flag is set the MRC will setup the following memory map for NUMA systems:

- All normal memory from socket 0 interleaved together
- All non-volatile memory(NVDIMM or BBU memory) from socket 0 interleaved together
- All normal memory from socket 1 interleaved together
- All non-volatile memory (NVDIMM or BBU memory) from socket 1 interleaved together

For example, if the system contains 4 Normal 4GB DIMMs on Socket 0, 2 4GB NVDIMMs on Socket 0, 8 Normal 4GB DIMMs on Socket 1 and 6 4GB NVDIMMs on Socket 1, the memory map will be the following assuming TOLM is set to 2GB

- 0-2GB: Socket 0 Normal memory
- 4-18GB: Socket 0 Normal memory
- 18-26GB: Socket 0 non-volatile memory
- 26-58GB: Socket 1 Normal memory
- 58-82GB: Socket 1 non-volatile memory

If NUMA is disabled the memory map will be:

- All non-NVDIMM DIMMs from socket both sockets interleaved together
- All NVDIMM DIMMs from both sockets interleaved together

